

# Matrix factorization

Suraj Kumar

Inria & ENS Lyon

Email: *suraj.kumar@inria.fr*

CR12: September 2023

<https://surakuma.github.io/courses/daamtc.html>

# Matrix factorizations

- Useful to solve systems of linear equations  $Ax = b$
- Popular factorizations
  - LU factorization
  - QR factorization
  - Singular value decomposition

# Important definitions

## Vector norm for $x \in \mathbb{R}^n$

The Euclidean norm of  $x$  is represented as  $\|x\|$  or  $\|x\|_2$  and defined as

$$\|x\| = \sqrt{\sum_{i=1}^n x_i^2}$$

## Matrix norm for $A \in \mathbb{R}^{n \times n}$

$$\text{Frobenius norm, } \|A\|_F = \sqrt{\sum_{j=1}^n \sum_{i=1}^n A_{ij}^2} = \sqrt{\text{trace}(AA^T)}$$

Spectral norm,  $\|A\|_2 = \text{largest singular value of } A$

## Orthogonal matrix

An orthogonal matrix  $Q$  satisfies  $Q^T Q = Q Q^T = I$  (the identity matrix)

- $Q$ 's rows are orthogonal to each other and have unit norm
- $Q$ 's columns are orthogonal to each other and have unit norm

# Table of Contents

1 Singular value decomposition

2 LU factorization

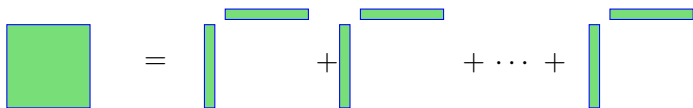
3 QR factorization

# Singular Value Decomposition (SVD)

- It decomposes a matrix  $A \in \mathbb{R}^{m \times n}$  to the form  $U\Sigma V^T$ 
  - $U$  is an  $m \times m$  orthogonal matrix
  - $V$  is an  $n \times n$  orthogonal matrix
  - $\Sigma$  is an  $m \times n$  rectangular diagonal matrix
- The diagonal entries  $\sigma_i = \Sigma_{ii}$  of  $\Sigma$  are called singular values
  - $\sigma_i \geq 0$  and  $\sigma_1 \leq \sigma_2 \leq \dots \leq \sigma_{\min(m,n)}$
- Columns of  $U$  and  $V$  are known as left and right singular vectors respectively
- If  $u_i, v_i$  are the  $i$ th vector of  $U$  and  $V$ , then  $A = \sum_{i=1}^{\min(m,n)} \sigma_i u_i v_i^T$
- The largest  $r$  such that  $\sigma_r \neq 0$  is called the rank of the matrix

# SVD and rank of a matrix

- SVD represents a matrix as the sum of  $r$  rank one matrices



- $\|A\|_F^2 = \sum_{i=1}^{\min(m,n)} \sigma_i^2 = \sum_{i=1}^r \sigma_i^2$
- If  $r' \leq r$  and  $\tilde{A} = \sum_{i=1}^{r'} \sigma_i u_i v_i^T$ , then

$$\|A - \tilde{A}\|_F^2 = \sum_{i=r'+1}^{\min(m,n)} \sigma_i^2 = \sum_{i=r'+1}^r \sigma_i^2$$

- Useful for compression, dimension reduction and low-rank approximation
- Expensive to compute and hard to parallelize

# Table of Contents

- 1 Singular value decomposition
- 2 LU factorization
- 3 QR factorization

# Algebra of LU factorization with an example

Given the matrix  $A = \begin{pmatrix} 2 & 6 & 5 \\ 4 & 15 & 11 \\ 6 & 30 & 23 \end{pmatrix}$

- Let  $L_1 = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ -3 & 0 & 1 \end{pmatrix}$ ,  $L_1 A = \begin{pmatrix} 2 & 6 & 5 \\ 0 & 3 & 1 \\ 0 & 12 & 8 \end{pmatrix}$

- Let  $L_2 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -4 & 1 \end{pmatrix}$ ,  $L_2 L_1 A = \begin{pmatrix} 2 & 6 & 5 \\ 0 & 3 & 1 \\ 0 & 0 & 4 \end{pmatrix}$

- Let  $U = \begin{pmatrix} 2 & 6 & 5 \\ 0 & 3 & 1 \\ 0 & 0 & 4 \end{pmatrix}$ ,  $L_2 L_1 A = U$



# Algebra of LU factorization

$$L_2 L_1 A = U \implies A = (L_2 L_1)^{-1} U = L_1^{-1} L_2^{-1} U$$

$$L_1 = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ -3 & 0 & 1 \end{pmatrix}, L_1^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 0 & 1 \end{pmatrix}, L_2 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -4 & 1 \end{pmatrix}, L_2^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 4 & 1 \end{pmatrix}$$

$$L_1^{-1} L_2^{-1} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 4 & 1 \end{pmatrix}$$

$$A = \begin{pmatrix} 2 & 6 & 5 \\ 4 & 15 & 11 \\ 6 & 30 & 23 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 3 & 4 & 1 \end{pmatrix} \begin{pmatrix} 2 & 6 & 5 \\ 0 & 3 & 1 \\ 0 & 0 & 4 \end{pmatrix} = LU, \text{ where } L = L_1^{-1} L_2^{-1}$$

# The need of pivoting (or row exchanges): $PA = LU$

- To avoid division by 0 or small diagonal elements (for stability)
- $A = \begin{pmatrix} 0 & 2 & 4 \\ 3 & 1 & 2 \\ 6 & 8 & 7 \end{pmatrix}$  has an LU factorization if we permute the rows of the matrix  $A$

$$PA = \begin{pmatrix} 3 & 1 & 2 \\ 0 & 2 & 4 \\ 6 & 8 & 7 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 2 & 3 & 1 \end{pmatrix} \begin{pmatrix} 3 & 1 & 2 \\ 0 & 2 & 4 \\ 0 & 0 & -9 \end{pmatrix}$$

$$\text{Here } P = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

# Communication lower bounds

- Matrix multiplication lower bounds apply to LU factorization using reduction [Ballard et. al., 09]

$$\begin{pmatrix} I & & -B \\ A & I & \\ & & I \end{pmatrix} = \begin{pmatrix} I & & \\ A & I & \\ & & I \end{pmatrix} \begin{pmatrix} I & & -B \\ & I & AB \\ & & I \end{pmatrix}$$


## Lower bounds

- Sequential lower bound on bandwidth =  $\Omega(\frac{n^3}{\sqrt{M}})$
- Memory-dependent parallel lower bound on bandwidth =  $\Omega(\frac{n^3}{P\sqrt{M}})$
- Memory-independent parallel lower bound on bandwidth =  $\Omega\left(\frac{n^3}{P^{\frac{2}{3}}}\right)$

# LU factorization

LU factorization (Gaussian elimination):

- Convert a matrix  $A$  into product  $L \times U$
- $L$  is lower triangular with diagonal 1
- $U$  is upper triangular
- $L$  and  $U$  stored in place with  $A$



## LU Algorithm

For  $k = 1 \dots n - 1$ :

- For  $i = k + 1 \dots n$ ,  
 $A_{i,k} \leftarrow A_{i,k} / A_{k,k}$  (column/panel preparation)
- For  $i = k + 1 \dots n$ ,  
For  $j = k + 1 \dots n$ ,  
 $A_{i,j} \leftarrow A_{i,j} - A_{i,k} A_{k,j}$  (update)

# Block LU factorization

## Partition of a $n \times n$ matrix $A$

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}$$

Here  $A_{11}$  is of size  $b \times b$ ,  $A_{21}$  is of size  $(n - b) \times b$ ,  $A_{12}$  is of size  $b \times (n - b)$  and  $A_{22}$  is of size  $(n - b) \times (n - b)$ .

## Structure of LU factorization algorithm

- The first iteration computes the factorization:

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} = \begin{pmatrix} L_{11} & \\ L_{21} & I_{n-b} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} \\ & A' \end{pmatrix}$$

- The algorithm continues recursively on the trailing matrix  $A'$ .

# Block LU factorization

- 1 Compute the LU factorization of the first block column

$$\begin{pmatrix} A_{11} \\ A_{21} \end{pmatrix} = \begin{pmatrix} L_{11} \\ L_{21} \end{pmatrix} U_{11}$$

- 2 Solve the triangular system

$$L_{11} U_{12} = A_{12}$$

- 3 Update the trailing matrix

$$A' = A_{22} - L_{21} U_{12}$$

- 4 Compute recursively the block LU factorization of  $A'$

# Table of Contents

- 1 Singular value decomposition
- 2 LU factorization
- 3 QR factorization

# Terminology related to QR factorization

An **orthogonal** matrix  $Q$  satisfies  $Q^T Q = Q Q^T = I$  (the identity matrix)

- $Q$  must be square
- $Q$ 's rows are orthogonal to each other and have unit norm
- $Q$ 's columns are orthogonal to each other and have unit norm



# Terminology related to QR factorization

An **orthogonal** matrix  $Q$  satisfies  $Q^T Q = Q Q^T = I$  (the identity matrix)

- $Q$  must be square
- $Q$ 's rows are orthogonal to each other and have unit norm
- $Q$ 's columns are orthogonal to each other and have unit norm

A matrix  $U$  has **orthonormal columns** if  $U^T U = I$  (the identity matrix)

- $U$ 's columns are orthogonal to each other and have unit norm
- $U$  can have more rows than columns, in which case  $U U^T \neq I$

# Terminology related to QR factorization

An **orthogonal** matrix  $Q$  satisfies  $Q^T Q = Q Q^T = I$  (the identity matrix)

- $Q$  must be square
- $Q$ 's rows are orthogonal to each other and have unit norm
- $Q$ 's columns are orthogonal to each other and have unit norm

A matrix  $U$  has **orthonormal columns** if  $U^T U = I$  (the identity matrix)

- $U$ 's columns are orthogonal to each other and have unit norm
- $U$  can have more rows than columns, in which case  $U U^T \neq I$

Given a matrix  $A$ , we can **orthogonalize** its columns by finding a matrix  $Q$  such that

- $Q$ 's columns span the same space as  $A$ 's columns
- $Q$  has orthonormal columns
- there exists a nonsingular matrix  $Z$  such that  $A = QZ$

The **QR factorization** is a fundamental matrix factorization:

$$A = QR = [\hat{Q} \quad \tilde{Q}] \begin{bmatrix} \hat{R} \\ 0 \end{bmatrix} = \hat{Q} \hat{R}$$

- if  $A$  is  $m \times n$ ,  $m \geq n$ , then  $Q$  is  $m \times m$ ,  $R$  is  $m \times n$ ,  $\hat{Q}$  is  $m \times n$ , and  $\hat{R}$  is  $n \times n$
- $Q$  is orthogonal,  $\hat{Q}$  has orthonormal columns, and  $R$  is upper triangular
- $\hat{Q}$  is an orthogonalization of  $A$

# Classical algorithms for QR factorization

## ① Gram-Schmidt process

- intuitive: each vector is orthogonalized against previous ones by subtracting out components of the vector in previous directions
- has numerical problems (vectors aren't always numerically orthonormal)
- two variants “classical” and “modified” are mathematically identical

## ② Householder QR

- uses orthogonal matrices to transform input to triangular form
- numerically stable

## Classical Gram-Schmidt (CGS) process

**Require:**  $A = [x_1 \ x_2 \ \cdots \ x_n]$

**for**  $i = 1$  to  $n$  **do**

$v_i = x_i$

**for**  $j = 1$  to  $i - 1$  **do**

$r_{ji} = q_j^T x_i$   $\triangleright$  compute size of projection of  $i$ th col of  $A$  onto  $q_j$

$v_i = v_i - r_{ji}q_j$   $\triangleright$  remove this component from vector  $v_i$

**end for**

$r_{ii} = \|v_i\|_2$

$q_i = v_i / r_{ii}$   $\triangleright$  normalize vector

**end for**

**Ensure:**  $Q = [q_1 \ q_2 \ \cdots \ q_n]$  has orthonormal columns

**Ensure:**  $R$  is upper triangular and  $A = QR$

## Modified Gram-Schmidt (MGS) process

**Require:**  $A = [x_1 \ x_2 \ \cdots \ x_n]$

**for**  $i = 1$  to  $n$  **do**

$v_i = x_i$

**for**  $j = 1$  to  $i - 1$  **do**

$r_{ji} = q_j^T v_i$   $\triangleright$  compute size of projection of **current vector** onto  $q_j$

$v_i = v_i - r_{ji} q_j$   $\triangleright$  remove this component from vector  $v_i$

**end for**

$r_{ii} = \|v_i\|_2$

$q_i = v_i / r_{ii}$   $\triangleright$  normalize vector

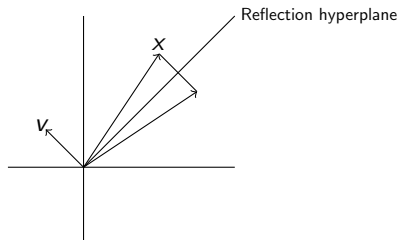
**end for**

**Ensure:**  $Q = [q_1 \ q_2 \ \cdots \ q_n]$  has orthonormal columns

**Ensure:**  $R$  is upper triangular and  $A = QR$

# Householder transformation

- $v$  is a unit vector
- The reflection hyperplane can be defined by its normal vector  $v$
- $(I - 2vv^T)x$  is the reflection of point  $x$  with the hyperplane



- $P = I - 2vv^T$  matrix is known as the Householder matrix
- $P$  is symmetric and orthogonal,  $P^2 = I$

# Main idea of Householder QR factorization

Look for a Householder matrix that annihilates the elements of a vector  $x$ , except first one:

$$Px = y, \|x\|_2 = \|y\|_2, y = \sigma e_1, \sigma = \pm \|x\|_2$$

The choice of sign is made to avoid cancellation or small numerical values while computing  $v_1 = x_1 - \sigma$ . Here  $v_1, x_1$  are the first elements of vectors  $v, x$  respectively.

$$v = x - y = x - \sigma e_1$$

$$\sigma = -\text{sign}(x_1) \|x\|_2, v = x - \sigma e_1$$

$$u = \frac{v}{\|v\|_2}, P = I - 2uu^T$$



# Householder QR algorithm

Given vector  $x$ , a **Householder transformation**  $I - 2uu^T$  maps  $x$  to  $\sigma e_1$

- $u$  is called the **Householder vector**

**Require:**  $A = \begin{bmatrix} x_1 & x_2 & \cdots & x_n \end{bmatrix}$

**for**  $i = 1$  to  $n$  **do**

    Compute Householder vector  $u_i$  from  $x_i$

$A = (I - 2u_i u_i^T)A$  ▷ apply Householder transformation

**end for**

$R = A$

**Ensure:**  $U = \begin{bmatrix} u_1 & u_2 & \cdots & u_n \end{bmatrix}$  is lower triangular

**Ensure:**  $R$  is upper triangular and  $A = (I - 2u_1 u_1^T) \cdots (I - 2u_n u_n^T)R$

# Householder QR computational complexity

Let  $A \in \mathbb{R}^{m \times n}$ , we count the number of operation to update  $A$  ( $A = (I - 2u_i u_i^T)A = A - 2u_i u_i^T A$ ) in each iteration  $i$ .

## Operations per iteration

- Dot product  $w = u_i^T A(i : m, i : n) : 2(m - i)(n - i)$
- Outer product  $u_i w : (m - i)(n - i)$
- Subtraction  $A(i : m, i : n) = A(i : m, i : n) - 2u_i w : (m - i)(n - i)$

The number of operations to multiply 2 with  $w$  is  $(n - i)$ , however it is a lower order term. Hence we do not consider it explicitly.

## Operations in Householder QR factorization

$$\begin{aligned}\sum_{i=1}^n &= 4(m - i)(n - i) = 4 \sum_{i=1}^n = 4(mn - (m + n)i + i^2) \\ &\approx 4mn^2 - 4(m + n)\frac{n^2}{2} + 4\frac{n^3}{3} = 2mn^2 - 2\frac{n^3}{3}\end{aligned}$$

# QR factorization

- $Q$  can be stored in compact representation
- Structure of block QR algorithm is similar to the block LU algorithm
- Matrix communication lower bounds are also valid for the Householder/CGS/MGS QR factorization

# QR factorization algorithms

Algorithm	# flops	# words	stability
CGS	$2mn^2 + O(n^3)$	$O(mn^2)$	Bad
MGS		$O(mn^2)$	Okay
HouseholderQR		$O(mn^2)$	Good

# Table of Contents

- 1 Singular value decomposition
- 2 LU factorization
- 3 QR factorization
  - Tall Skinny QR (TSQR) factorization

# TSQR: QR factorization of a tall skinny matrix

QR factorization of a  $m \times n$  matrix with  $m \gg n$


The goal and process of Householder QR:

- annihilate entries below diagonal to obtain upper triangular form
- work column-by-column, left-to-right

Tall-Skinny QR idea (Demmel, Grigori, Hoemmen, Langou '12):

- change the order of annihilation to minimize communication
- work row-by-row, top to bottom

# Algebra of TSQR



$$A = \begin{pmatrix} A_1 \\ A_2 \\ A_3 \\ A_4 \end{pmatrix} = \begin{pmatrix} Q_1^{(2)} R_1^{(2)} \\ Q_2^{(2)} R_2^{(2)} \\ Q_3^{(2)} R_3^{(2)} \\ Q_4^{(2)} R_4^{(2)} \end{pmatrix} = \begin{pmatrix} Q_1^{(2)} & & & \\ & Q_2^{(2)} & & \\ & & Q_3^{(2)} & \\ & & & Q_4^{(2)} \end{pmatrix} \begin{pmatrix} R_1^{(2)} \\ R_2^{(2)} \\ R_3^{(2)} \\ R_4^{(2)} \end{pmatrix}$$

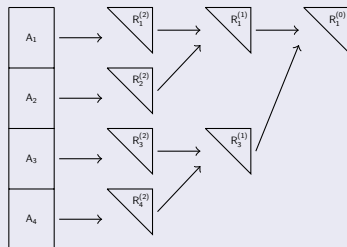
$$\begin{pmatrix} R_1^{(2)} \\ R_2^{(2)} \\ R_3^{(2)} \\ R_4^{(2)} \end{pmatrix} = \begin{pmatrix} Q_1^{(1)} R_1^{(1)} \\ Q_2^{(1)} R_2^{(1)} \end{pmatrix} = \begin{pmatrix} Q_1^{(1)} & \\ & Q_2^{(1)} \end{pmatrix} \begin{pmatrix} R_1^{(1)} \\ R_2^{(1)} \end{pmatrix}, \quad \begin{pmatrix} R_1^{(1)} \\ R_2^{(1)} \end{pmatrix} = Q_1^{(0)} R_1^{(0)}$$

$Q$  is represented implicitly as a product.

# Flexibility of TSQR

## Parallel TSQR

- Assuming block row layout on  $P$  processors
- Communication cost is that of binomial-tree reduction:  
 $\beta \cdot O(n^2 \log P) + \alpha \cdot O(\log P)$



## Sequential TSQR

- Assuming cache size is  $\Omega(n^2)$
- It streams through matrix once achieving  $O(mn)$  amount of data transfers

