

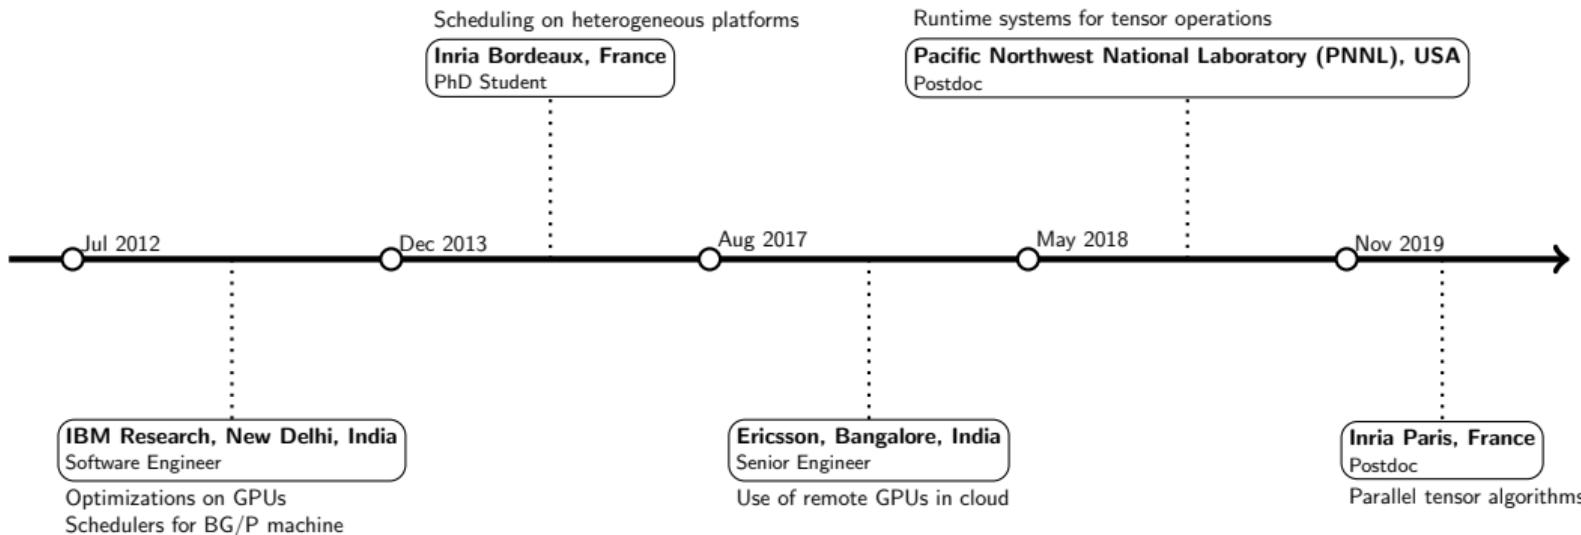
Scalable Tensor Algorithms for Modern Computing Systems

Suraj KUMAR

CNRS LIP/LaBRI Applicant

March 15, 2022

Resume in timeline

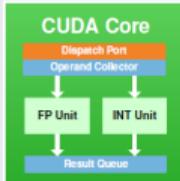


Advisors/Close collaborators

- PhD supervisors at Inria Bordeaux, France: Olivier Beaumont, Lionel Eyraud-Dubois, Samuel Thibault, Emmanuel Agullo
- Postdoc supervisors/Close collaborators: Marcin Zalewski (PNNL, USA), Sriram Krishnamoorthy (PNNL, USA), Laura Grigori (Inria Paris, France), Grey Ballard (Wake Forest University, USA)

Past research experience

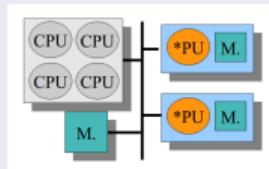
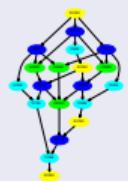
GPU algorithms



IBM Research, India, 2013

- Stencil computations

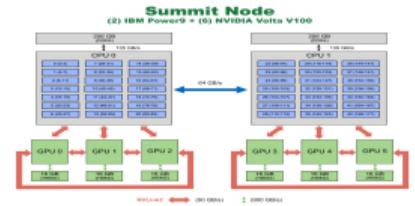
Scheduling on heterogeneous platforms



PhD thesis, Inria Bordeaux, France, 2017

- Dense linear algebra computations

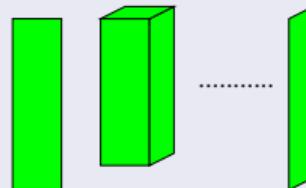
Fast molecular simulations



PNNL, USA, 2019

- Tensor computations, TAMM library

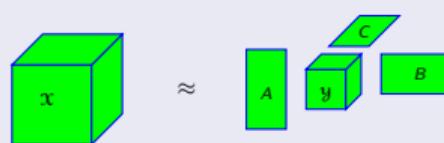
Parallel tensor approximations



Inria Paris, France, 2019

- ■ ■ Tensor-train decomposition

Communication optimal algorithms



Inria Paris, France, 2019

- ■ ■ Multiple Tensor-times-matrix computations (to obtain \mathcal{Y})

Expertise: **Tensor computations, Parallel algorithms, Communication costs, Low-rank approximations**

Tensors and their uses

- **Neuroscience:** Neuron \times Time \times Trial
 - **Media:** User \times Movie \times Time
 - **Ecommerce:** User \times Product \times Time
 - **Social-Network:** Person \times Person \times Time \times Type
-
- High dimensional tensors: Neural network, Molecular simulation, Quantum computing
 - People work with low dimensional structure (decomposition) of the tensors



Canonical decomposition

$$\text{Tensor} = \text{Factor}_1 + \text{Factor}_2 + \dots + \text{Factor}_n$$

Tucker decomposition

$$\text{Tensor} = \text{Core} \times \text{Factor}_1 \times \text{Factor}_2 \times \dots \times \text{Factor}_n$$

Tensor-train decomposition

$$\text{Tensor} = \text{Column}_1 \otimes \text{Column}_2 \otimes \dots \otimes \text{Column}_n$$

Importance of communication

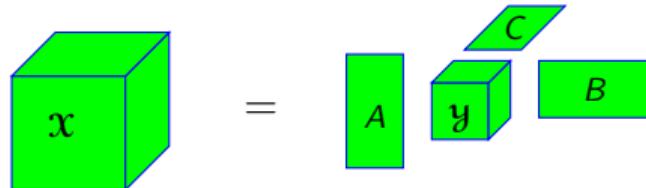
- Gaps between computation and communication costs growing exponentially

	time-per-operation	Network-bandwidth	Network-latency
Annual improvements	59 %	26 %	15 %

Source: Getting up to speed: The future of supercomputing

Goal : Scalable and communication optimal algorithms for tensor computations

Higher-order SVD (HOSVD) to compute Tucker decomposition



Algorithm 1 3-dimensional HOSVD Algorithm(\mathcal{X})

- 1: Obtain factor matrices A, B and C from the matrix representations of the input tensor \mathcal{X}
 - 2: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$
 - 3: Return \mathcal{Y}, A, B, C
-

- \mathcal{X}, \mathcal{Y} : 3-dimensional input and output tensors (or arrays) & A, B, C : matrices
- \times_i : tensor contraction along the i th dimension (similar to matrix multiplication)
- Multiple Tensor-Times-Matrix (Multi-TTM) computation: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$

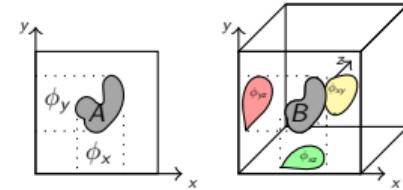
- 1 For Matrix Matrix Multiplications
- 2 For Multi-TTM Computations

Settings

- P number of processors
- Each processor performs (asymptotically) equal amount of operations
- One copy of data is in the system
 - $1/P$ th amount of inputs (before the computation) and output (after the computation) on each processor
- Focus on bandwidth cost (volume of data transfers)
- Submitted this work to SPAA 2022 conference
- Joint work with L. Grigori, G. Ballard, K. Rouse, H. Al Daas

Approach to obtain communication lower bounds

- Loomis-Whitney inequality: for $d - 1$ dimensional projections
 - For the 2d object A , $\phi_x \phi_y \geq \text{Area}(A)$
 - For the 3d object B , $(\phi_{xy} \phi_{yz} \phi_{xz})^{\frac{1}{2}} \geq \text{Volume}(B)$
- Hölder-Brascamp-Lieb (HBL) inequality – generalization for arbitrary dimensional projections
 - Provide exponent for each projection



Constraints for parallel load balanced matrix matrix multiplication

- $C = AB$ with $A \in \mathbb{R}^{n_1 \times n_2}$, $B \in \mathbb{R}^{n_2 \times n_3}$, and $C \in \mathbb{R}^{n_1 \times n_3}$
for $i = 1:n_1$, for $k = 1:n_2$, for $j = 1:n_3$
$$C[i][j] += A[i][k] * B[k][j]$$
- ϕ_A, ϕ_B, ϕ_C : projections of computations on arrays A, B, C
- From Loomis-Whitney/HBL inequality: $\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \text{number of multiplications per processor} = \frac{n_1 n_2 n_3}{P}$
- Our contributions: $\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$, $\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$, $\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$

Optimization problem and Communication lower bounds

Minimize $\phi_A + \phi_B + \phi_C$ s.t.

$$\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \frac{n_1 n_2 n_3}{P}$$

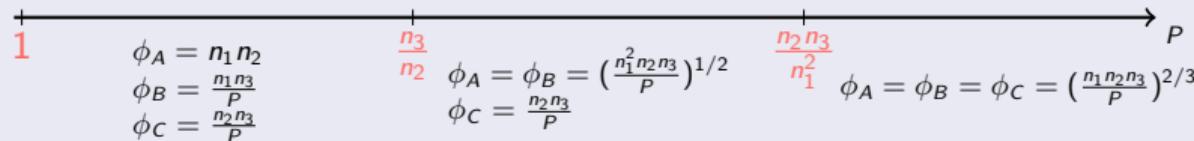
$$\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$$

$$\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$$

$$\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$$

Amount of array accesses = $\phi_A + \phi_B + \phi_C$

- Estimate the solution and prove optimality using all Karush–Kuhn–Tucker conditions are satisfied
- For $n_1 \leq n_2 \leq n_3$,



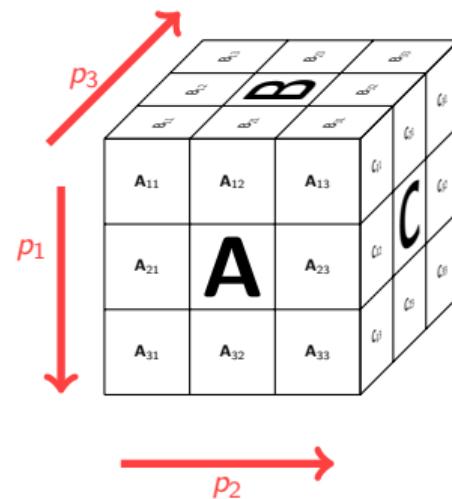
- Communication lower bound = $\phi_A + \phi_B + \phi_C - \text{data owned by the processor} = \phi_A + \phi_B + \phi_C - \frac{n_1 n_2 + n_2 n_3 + n_1 n_3}{P}$

Design of communication optimal algorithms for $C = AB$

Arrangements of 8 processors



- P is organized into $p_1 \times p_2 \times p_3$ logical grid
- Select p_1, p_2 and p_3 based on the communication lower bounds
- Gather A on the set of processors along each slice of p_3
- Gather B on the set of processors along each slice of p_1
- Perform local computation
- Perform reduce operation along p_2 to obtain C



3-dimensional Multi-TTM ($\mathcal{Y} = \mathcal{X} \times_1 \mathbf{A}^{(1)T} \times_2 \mathbf{A}^{(2)T} \times_3 \mathbf{A}^{(3)T}$)

- TTM-in-Sequence approach (used in Tucker-MPI)

- $\mathcal{Y} = \left(\left(\mathcal{X} \times_1 \mathbf{A}^{(1)T} \right) \times_2 \mathbf{A}^{(2)T} \right) \times_3 \mathbf{A}^{(3)T}$

All-at-Once definition with $\mathcal{X} \in \mathbb{R}^{n_1 \times n_2 \times n_3}$, $\mathcal{Y} \in \mathbb{R}^{r_1 \times r_2 \times r_3}$, $\mathbf{A}^{(k)} \in \mathbb{R}^{n_k \times r_k}$ (our contribution)

for $n'_1 = 1:n_1$, for $n'_2 = 1:n_2$, for $n'_3 = 1:n_3$

for $r'_1 = 1:r_1$, for $r'_2 = 1:r_2$, for $r'_3 = 1:r_3$

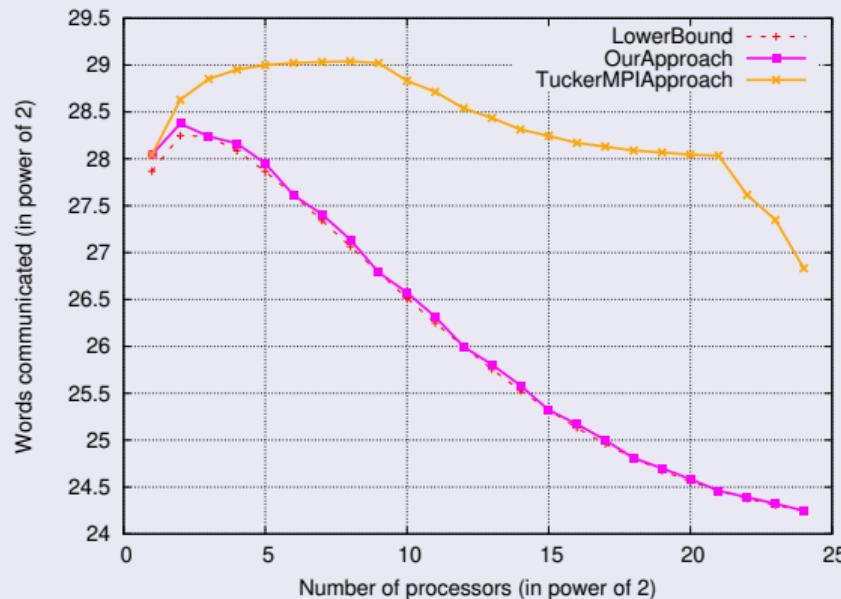
$$\mathcal{Y}(r'_1, r'_2, r'_3) = \mathcal{Y}(r'_1, r'_2, r'_3)$$

$$+ \left(\mathcal{X}(n'_1, n'_2, n'_3) * \mathbf{A}^{(1)}(n'_1, r'_1) * \mathbf{A}^{(2)}(n'_2, r'_2) * \mathbf{A}^{(3)}(n'_3, r'_3) \right)$$

- Applied our framework to compute lower bounds
- Designed similar communication optimal algorithms (though 6-dimensional)

Simulated performance comparison

$$n_1 = n_2 = n_3 = 2^{20}, r_1 = r_2 = r_3 = 2^8$$



- Typical scenarios in data compression problems
- Our approach communicates much less than the state-of-the-art approach (TuckerMPI)

Project: Scalable Tensor Algorithms for Modern Computing Systems

- 1 Design of Scalable Communication Optimal Algorithms for Tensors (Main Focus)
- 2 Extension of Existing Approaches/Algorithms (Short/Mid Term Research Plans)
- 3 Exploratory Topics (Mid/Long Term Research Plans)

Scalable algorithms for popular tensor operations

- Determine the communication lower bounds for tensor decompositions
- Analyse the popular decomposition algorithms and communications performed by them
- Propose new scalable communication optimal algorithms
 - If possible design tiles/tasks based algorithms
- Implement the proposed algorithms
- Same for manipulation operations of popular tensor representations
- Create a tensor library

Strassen's concepts to tensors

Matrix multiplication of $n \times n$ square matrices

- Complexity of traditional matrix multiplication is $\mathcal{O}(n^3)$
- Strassen's matrix multiplication
 - Expressed matrix multiplication as a tensor computation
 - Canonical rank of the tensor determines the complexity of the computation
 - Complexity is $\mathcal{O}(n^{2.81})$
- Plan to extend Strassen's concepts to tensor contractions

Contraction of a 3-dimensional tensor with a matrix

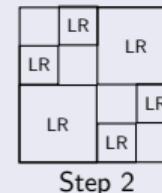
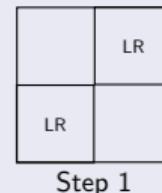
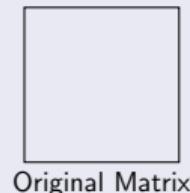
```
for  $i_1 = 1 : n$  do
    for  $i_2 = 1 : n$  do
        for  $i_3 = 1 : n$  do
            for  $j_2 = 1 : n$  do
                 $G(i_1, i_2, j_2) = G(i_1, i_2, j_2) + A(i_1, i_2, i_3) * B(i_3, j_2)$ 
            end for
        end for
    end for
end for
```

- Total $\mathcal{O}(n^4)$ operations
- Apply Strassen's algorithm for each i_1 , total $\mathcal{O}(n^{3.81})$ operations
- Expressing as a canonical decomposition of $8 \times 8 \times 4$ tensor can further reduce the number of operations (first try)

Hierarchical matrix concepts to tensors

Hierarchical matrices

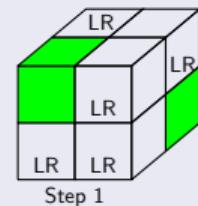
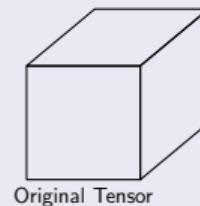
- Data sparse approximation of non-sparse matrices



LR: low rank block

Tensors

- $f(i, j, k) = \frac{1}{|i-j|+|j-k|+|k-i|}$
- Value is small if difference of any pair is large
- Formalize and evaluate this approach for tensors



High dimensional tensor representations

Notation: Tensors are denoted by solid shapes and number of lines denote the dimensions of the tensors. Connecting two lines implies summation (or contraction) over the connected dimensions.

- Adding tensors and applying an operator in Tensor-train representation

$$\begin{array}{c} \text{rank}=r \\ \text{---} \\ \text{rank}=r \end{array} + \begin{array}{c} \text{rank}=s \\ \text{---} \\ \text{rank}=s \end{array} = \begin{array}{c} \text{rank}=r+s \\ \text{---} \\ \text{rank}=r+s \end{array}$$
$$\begin{array}{c} \text{rank}=r \\ \text{---} \\ \text{rank}=r \end{array} = \begin{array}{c} \text{rank}=s \\ \text{---} \\ \text{rank}=s \end{array} = \begin{array}{c} \text{rank}=r*s \\ \text{---} \\ \text{rank}=r*s \end{array}$$

- Requires a truncation process which iterates over cores one by one – not suitable to work in parallel

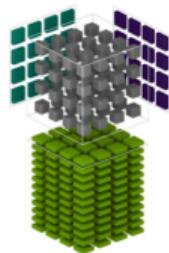
New tensor representations – suitable for parallelization

- Look at new representations in tree format



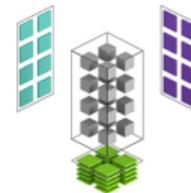
- Data will be stored at the leaf nodes

Architecture aware algorithms and Training of tensorized neural networks



$$D = \begin{pmatrix} K_1 & K_2 & K_3 & K_4 \\ K_5 & K_6 & K_7 & K_8 \\ K_9 & K_{10} & K_{11} & K_{12} \\ K_{13} & K_{14} & K_{15} & K_{16} \end{pmatrix} + \begin{pmatrix} K_1 & K_2 & K_3 & K_4 \\ K_5 & K_6 & K_7 & K_8 \\ K_9 & K_{10} & K_{11} & K_{12} \\ K_{13} & K_{14} & K_{15} & K_{16} \end{pmatrix}$$

NVIDIA A100 Tensor Core FP64



- Recent Nvidia GPUs have tensor cores to accelerate AI computations
- Design parallel algorithms which take architecture details into account
- Design parallel methods to train tensorized neural networks

Integration in the LIP/LaBRI laboratory

My research plans

- **Main focus:** scalable and communication optimal tensor algorithms
- **Short term plans:** Strassen's concepts to tensors, hierarchical matrix concepts to tensors
- **Long term plans:** new tensor representations, parallel training of tensorized neural networks

ROMA team (LIP laboratory)

- *Bora Ucar*: design of tensor compression and manipulation algorithms
- *Gregoire Pichon*: low-rank based methods
- *Anne Benoit, Loris Marchal, Yves Robert and Frederic Vivien*: scalability and scheduling aspects in the long term

SATANAS team (LaBRI laboratory)

- *Olivier Beaumont and Lionel Eyraud-Dubois*: parallel training of tensorized neural networks
- *Mathieu Faverge*: low-rank based methods
- *Abdou Guermouche, Samuel Thibault*: exploitation of maximum potential of HPC systems in the long term

Bringing additional skills

- High dimensional dense tensor computations, use of tensors in molecular simulations (for both teams)
- Communication lower bounds for linear algebra computations (for both teams)
- Scalable approaches for large HPC systems and GPU computations (only for the ROMA team)

Thank You!