

# Scalable Tensor Algorithms for Modern Computing Systems

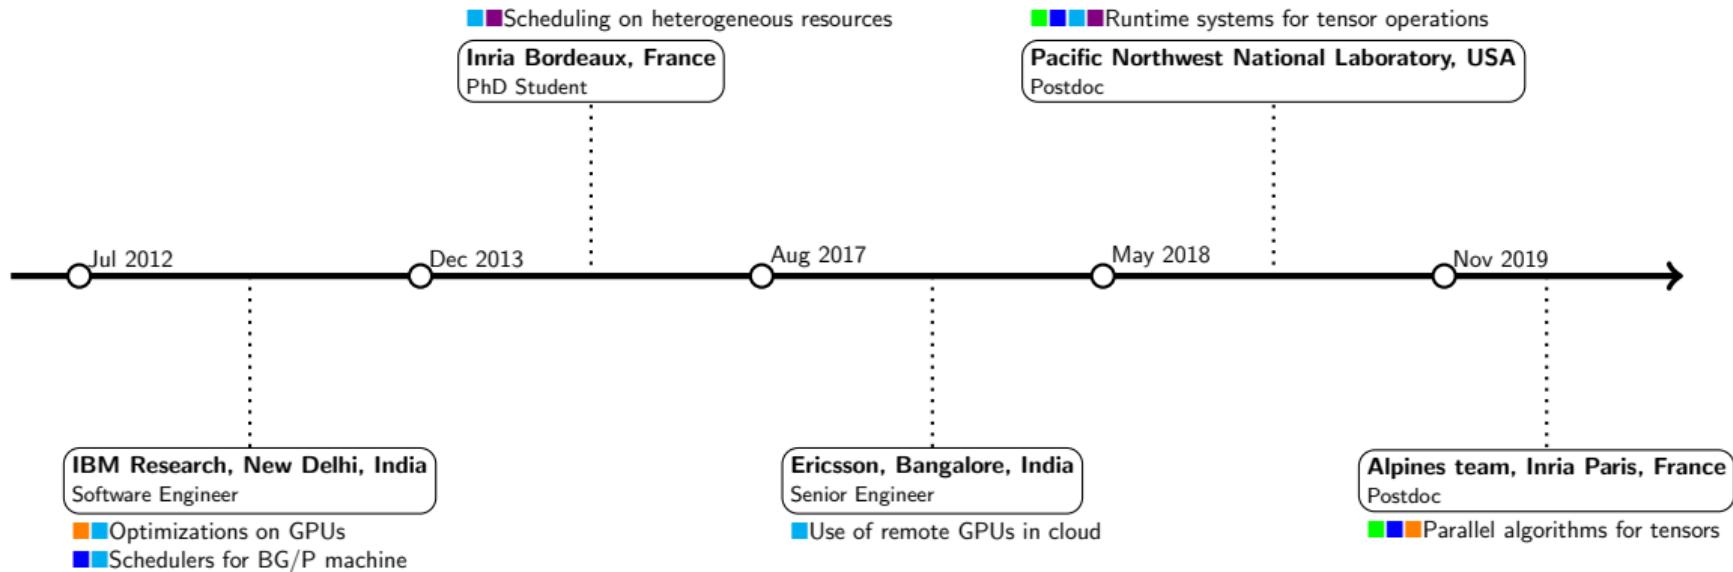
Suraj KUMAR

Inria ROMA Applicant

June 4, 2021

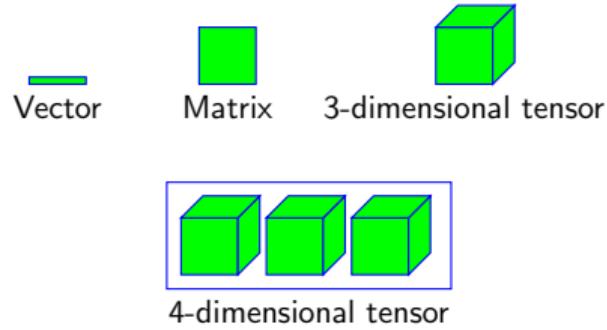
# Resume

Interests: **Tensors, Scalable Algorithms, Scheduling, Runtime Systems, Performance Optimizations**

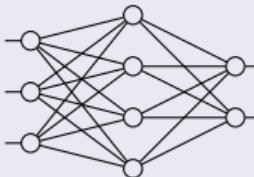


# Tensors are used in Several Domains

- **Neuroscience:** Neuron  $\times$  Time  $\times$  Trial
- **Transportation:** Pickup  $\times$  Dropoff  $\times$  Time
- **Media:** User  $\times$  Movie  $\times$  Time
- **Ecommerce:** User  $\times$  Product  $\times$  Time
- **Cyber-Traffic:** IP  $\times$  IP  $\times$  Port  $\times$  Time
- **Social-Network:** Person  $\times$  Person  $\times$  Time  $\times$  Interaction-Type



## High Dimensional Tensors

- **Neural Network:**  

- **Molecular Simulation:** To represent wave functions
- **Quantum Computing:** To represent qubit states

# Tensor Computations

- Memory and computation requirements are exponential in the number of dimensions
  - A simulation involving just 100 spatial orbitals manipulates a huge tensor with  $4^{100}$  elements
- People work with low dimensional structure (decomposition) of the tensors
  - A tensor is represented with smaller objects
  - Improves memory and computation requirements
- Most tensor decompositions rely on Singular Value Decomposition (SVD) of matrices
  - SVD represents a matrix as the sum of rank one matrices,  $A = U\Sigma V^T = \sum_i \Sigma(i; i)U_iV_i^T$

$$\begin{matrix} \text{[Large Green Square]} \end{matrix} = \begin{matrix} \text{[Small Green Square]} \end{matrix} \begin{matrix} \text{[Large Green Rectangle]} \end{matrix} = \begin{matrix} \text{[Thin Green Line]} \end{matrix} + \begin{matrix} \text{[Thin Green Line]} \end{matrix} + \cdots + \begin{matrix} \text{[Thin Green Line]} \end{matrix}$$

# Popular Tensor Decompositions (Higher Order Generalization of SVD)

- Canonical decomposition (Also known as Canonical Polyadic or CANDECOMP/PARAFAC)

$$\text{Tensor} = \text{rank-1 tensor} + \dots + \text{rank-1 tensor}$$

The diagram shows a large green cube representing the original tensor. To its right is an equals sign. Following the equals sign are three smaller green tensors, each consisting of two intersecting lines forming a cross shape, representing rank-1 tensors. A plus sign follows the first rank-1 tensor, and a ellipsis (three dots) follows the second rank-1 tensor, indicating they are summed together.

- Tucker decomposition

$$\text{Tensor} = \text{core tensor} \times \text{matrix}_1 \times \text{matrix}_2 \times \dots \times \text{matrix}_n$$

The diagram shows a large green cube representing the original tensor. To its right is an equals sign. Following the equals sign are three tensors: a small green rectangular block, a medium green rectangular block, and a large green rectangular block. The medium green block is positioned above the other two, suggesting it is the core tensor being multiplied by other factors.

- Tensor Train decomposition (equivalently known as Matrix Product States)

$$\mathbf{A} = \text{rank-1 tensor} \times \dots \times \text{rank-1 tensor}$$

The diagram shows a bold letter 'A' representing the tensor. To its right is an equals sign. Following the equals sign are four vertical green rectangles of increasing height, representing a sequence of rank-1 tensors being multiplied together.

Tensor notation: bold letters denote tensors, i.e.,  $\mathbf{A} \in \mathbb{R}^{n_1 \times \dots \times n_d}$  is a  $d$ -dimensional tensor.

# Previous/Ongoing Work

## 1 Parallel Tensor Train Algorithms

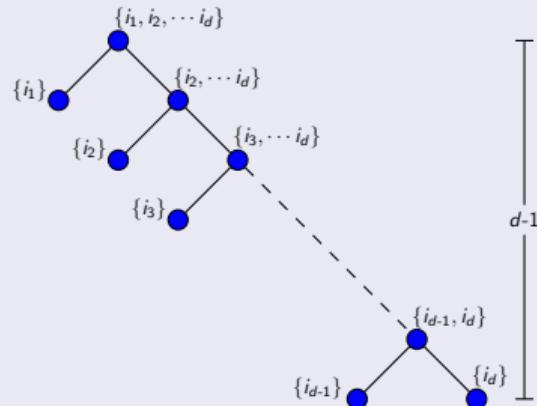
## 2 Scheduling on Heterogeneous Systems

- Scheduling of Dense Linear Algebra Kernels
- Communication Computation Overlap

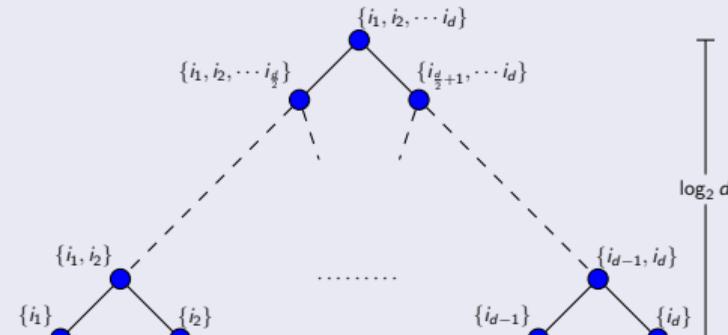
# Tensor Train algorithms and Separation of dimensions

- A sequential algorithm to compute Tensor Train decomposition exists [Oseledets, 2011]

## Sequential algorithm




## For better parallelization



- Can obtain better parallelism by expressing the operation in a balanced binary tree shape
  - Proposed a parallel algorithm based on this idea (joint work with L. Grigori, Inria Paris)


# Tensor Train approximation algorithms

## Sequential algorithm [Oseledets, 2011]



- Unfolding matrix: matricized representation of the tensor
- Perform truncated SVD of unfolding matrix  $A$ ,  $A = U\Sigma V^T + E_A$
- Work with  $\Sigma V^T$  on the right subtree

## Our parallel algorithm



- Perform truncated SVD of unfolding matrix  $A$ ,  $A = U\Sigma V^T + E_A$
- Find diagonal matrices  $X$ ,  $Y$ , and  $S$ , such that  $\Sigma = XSY$
- Call left (resp. right) subtree with  $UX$  (resp.  $YV^T$ )

Approach 1:  $X = I$ ,  $Y = \Sigma$ ,  $S = I$

Approach 2:  $X = Y = \Sigma^{\frac{1}{2}}$ ,  $S = I$

Approach 3:  $X = Y = \Sigma$ ,  $S = \Sigma^{-1}$

# Comparison of our approaches

- A 12-dimensional tensor with  $4^{12}$  elements (generated with a popular low rank function)
- prescribed accuracy =  $10^{-6}$
- Compr: compression ratio, NE: number of elements, AA: approximation accuracy

Metric	Sequential Algo	Parallel Algo		
		Approach 1	Approach 2	Approach 3
Compr	99.993	99.817	99.799	99.993
NE	1212	30632	33772	1212
AA	2.271e-07	3.629e-08	2.820e-08	2.265e-07

## SVD is expensive


- Good alternatives to SVD: QR factorization with column pivoting (QRCP), randomized SVD (RSVD)

Approach	Rank	Compr	NE	Sequential-AA	Approach3-AA
SVD	5	99.994	992	6.079e-06	6.079e-06
QRCP+SVD				1.016e-05	1.384e-05
RSVD				6.079e-06	6.079e-06
SVD	6	99.992	1376	1.323e-07	1.340e-07
QRCP+SVD				3.555e-07	5.737e-07
RSVD				1.322e-07	1.322e-07

# Performance Comparison

## Single core performance

- Number of computations for both RSVD algorithms =  $\mathcal{O}(n^d)$
- Approach3-SVD is very slow
- Approach3-RSVD is much faster




Parallel performance counts along the critical path on  $P$  processors

Algorithm	# Computations	Communications	# Messages
Sequential-RSVD	$\mathcal{O}\left(\frac{n^d}{P}\right)$	$\mathcal{O}\left(\frac{n^{d-1}}{\sqrt{P}} \log P\right)$	$\mathcal{O}(d \log P)$
Parallel-RSVD	$\mathcal{O}\left(\frac{n^d}{P}\right)$	$\mathcal{O}\left(\frac{n^{\frac{d}{2}}}{\sqrt{P}} \log P\right)$	$\mathcal{O}(\log d \log P)$

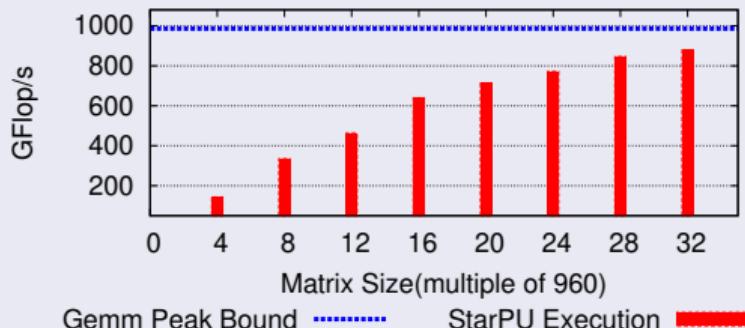
# Scheduling on Heterogeneous Systems

- Heterogeneous systems are common in High Performance Computing (HPC) (147 out of 500 in TOP500 list)
- Task based runtimes are a popular approach to exploit these systems



- Task based runtimes: StarPU, OmpSS, Legion, PaRSEC
- Application is represented as a graph of tasks (computations)
- E.g., Cholesky graph for  $4 \times 4$  tile matrix

## StarPU scheduler performance





- A platform with 9 CPUs and 3 GPUs
- Scheduler is based on popular heft strategy
- Goal: Enhance performance bounds and propose better scheduling strategies

Joint work with E. Agullo, O. Beaumont, L. Eyraud-Dubois, and S. Thibault during my PhD at Inria Bordeaux

# Our strategy and Performance comparison


## Trace for 12 X 12 tile matrix of Cholesky factorization



## Theoretical guarantees of our strategy and performance comparison

- Each resource selects the best suited task
- A fast resource restarts the blocking task



(#CPUs, #GPUs)	For a set of independent tasks	
	Approximation ratio	Worst case ex.
(1,1)	$\frac{1+\sqrt{5}}{2}$	$\frac{1+\sqrt{5}}{2}$
(m,n)	$2 + \sqrt{2} \approx 3.41$	$2 + \frac{2}{\sqrt{3}} \approx 3.15$



- Our upper bound is obtained by a linear program

# Minimizing impact of communications on Summit supercomputer

- Maximizing the overlap of communications and computations
- Implemented proposed approaches in Tensor Algebra for Manybody Methods (TAMM) library
- Molecular chemistry application (CCSD), Ubiqtin molecule, cc-pVDZ (737 basis functions, 220 nodes), aug-cc-pVDZ (1243 basis functions, 256 nodes)



Joint work with S. Krishnamoorthy and M. Zalewski during my postdoc at PNNL, USA

Fig source: <https://www.olcf.ornl.gov>

# Project: Scalable Tensor Algorithms for Modern Computing Systems

- 1 Design of Scalable Communication Optimal Algorithms for Tensors (Main Focus)
- 2 Extension of Existing Approaches/Algorithms (Short/Mid Term Research Plans)
- 3 Exploratory Topics (Mid/Long Term Research Plans)

## Scalable communication optimal algorithms for tensors

- Analyze existing algorithms
- Determine communication lower bounds
- Propose communication optimal algorithms
- Implement the proposed algorithms

Main focus

## Extension of existing approaches

- Strassen's concepts to tensors
- Concepts of hierarchical matrices to tensors
- Separation order of dimensions in tensor train

Short/Mid term plans


## Exploratory topics

- New tensor representations
- Architecture aware algorithms
- Randomization in tensors
- Factorizations of tensors

Mid/Long term plans

# Communication and its importance in HPC

- Running time of an algorithm depends on
  - Computations
    - Number of operations \* time-per-operation
  - Data movement
    - Volume of communication / Network-bandwidth
    - Number of messages \* Network-latency
- Gaps growing exponentially with time (Source: Getting up to speed: The future of supercomputing)



	time-per-operation	Network-bandwidth	Network-latency
Annual improvements	59 %	26 %	15 %


- Avoid communication to save time (and energy)

# Scalable algorithms for popular tensor operations

- Determine the communication lower bounds for tensor decompositions
- Analyse the popular decomposition algorithms and communications performed by them
- Propose new scalable communication optimal algorithms
  - If possible design tiles/tasks based algorithms
- Implement the proposed algorithms
  - Handle performance issues for homogeneous systems
    - Load balancing
    - Memory aware approaches
    - scheduling strategies
- Same for manipulation operations of popular tensor representations
- Extend implementation for heterogeneous systems (start with Nvidia GPUs based heterogeneous systems)
- Create a tensor library

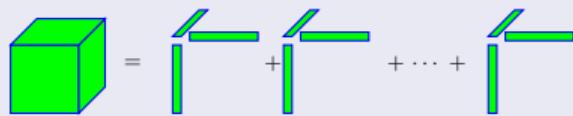
# Popular tensor decompositions

## Tucker decomposition




- Determine communication lower bounds for this operation
- Analyse communications performed by state of the art algorithms
- Propose and implement new scalable communication algorithms

## Canonical decomposition



- No deterministic algorithm to find the decomposition
- Analyse one iteration of the popular existing algorithms
- Propose and implement scalable algorithms for one iteration

## Tensor Train decomposition



- Determine communication lower bounds for this operation
- Analyse communication performed by popular algorithm
- Propose and implement new scalable communication algorithms

# Proving communication lower bounds for parallel computations

## How people did it for linear algebra operations?

- People obtain results for matrix multiplication operations
- Same lower bound apply to almost all direct linear algebra operations using reduction [Ballard et. al., 09] , for instance, bound for LU factorization

$$\begin{pmatrix} I & -B \\ A & I \\ \end{pmatrix} = \begin{pmatrix} I & \\ A & I \\ \end{pmatrix} \begin{pmatrix} I & -B \\ & AB \\ \end{pmatrix}$$

## Approach to compute lower bounds for tensor computations

Notation: Tensors are denoted by solid shapes and number of lines denote the dimensions of the tensors.  
Connecting two lines implies summation (or contraction) over the connected dimensions.

- Obtain bounds for basic tensor operations: Tensor times matrix (TTM), Multiple tensor times matrix (Multi-TTM), Tensor contraction



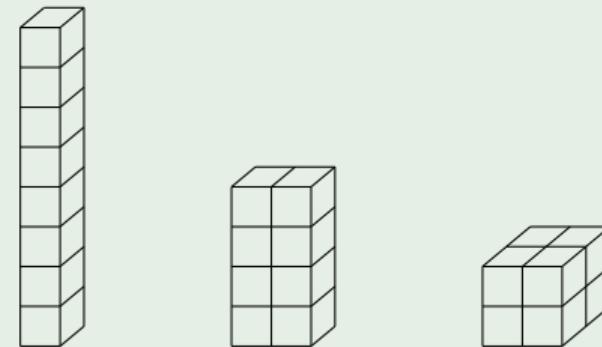
- Express decompositions and manipulations in terms of these basis operations

# Communication lower bounds on $P$ processors

- Recently started to work with L. Grigori (Inria Paris, France) and G. Ballard (Wake Forest University, USA)
- Revisited communication lower bounds for matrix multiplications

- Expressed existing approaches in suitable forms for tensors
- Lower bounds also instruct arrangement of processors in optimal algorithms
- Improved the constants in the existing ranges of  $P$  (Demmel et.al [IPDPS 2013])

Arrangements of 8 processors



- Plan to continue this collaboration to compute lower bounds for tensor computations

# Research Project

- 1 Design of Scalable Communication Optimal Algorithms for Tensors (Main Focus)
- 2 Extension of Existing Approaches/Algorithms (Short/Mid Term Research Plans)
- 3 Exploratory Topics (Mid/Long Term Research Plans)

# Strassen's concepts to tensors

## Matrix multiplication of $n \times n$ square matrices

- Complexity of traditional matrix multiplication is  $\mathcal{O}(n^3)$
- Strassen's matrix multiplication
  - Expressed matrix multiplication operation as a tensor operation
  - Canonical rank of the tensor determines the complexity of the operation
  - Complexity is  $\mathcal{O}(n^{2.81})$
- Plan to extend Strassen's concepts to tensor contractions

## Contraction of a 3-dimensional tensor with a matrix


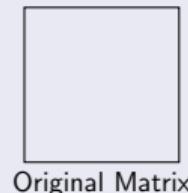
```
for i1 = 1 : n do
    for i2 = 1 : n do
        for i3 = 1 : n do
            for j2 = 1 : n do
                G(i1, i2, j2) = G(i1, i2, j2) + A(i1, i2, i3) * B(i3, j2)
            end for
        end for
    end for
end for
```

- Total  $\mathcal{O}(n^4)$  operations
- Apply Strassen's algorithm for each  $i_1$ , total  $\mathcal{O}(n^{3.81})$  operations
- Expressing as a canonical decomposition of  $8 \times 8 \times 4$  tensor can further reduce the number of operations


# Hierarchical Matrix concepts to Tensors

## Hierarchical Matrices

- Data sparse approximation of non-sparse matrices



Step 1




Step 2


*LR*: low rank block

## Tensors

- $f(i, j, k) = \frac{1}{|i-j|+|j-k|+|k-i|}$
- Value is small if difference of any pair is large
- Formalize and evaluate this approach for tensors



Original Tensor



Step 1

# Research Project

- 1 Design of Scalable Communication Optimal Algorithms for Tensors (Main Focus)
- 2 Extension of Existing Approaches/Algorithms (Short/Mid Term Research Plans)
- 3 Exploratory Topics (Mid/Long Term Research Plans)

# Tensor Representations for High Dimensional Tensors

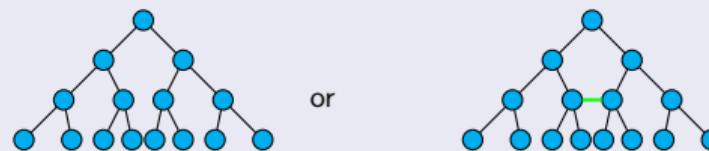
- Tensor Train is the popular representation to work with high dimensional tensors
- Adding tensors and applying an operator in this representation



- Requires a truncation process which iterates over cores one by one
- This representation is not much suited to work in parallel

## New Tensor Representations

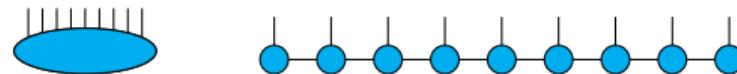
- Look at new representations in tree format – suitable for parallelization



- Data will be stored at the leaf nodes
- Internal nodes will help to manipulate tensors in parallel

# Randomization in Tensor Computations

- Randomized SVD and UTV factorization are now well established
- Apply randomization to tensors
- Perform factorizations of tensors
  - For example: QR like factorization of a tensor



# Integration in the ROMA team

Anne Benoit, Loris Marchal, Gregoire Pichon,  
Yves Robert, Bora Ucar, and Frederic Vivien

Team

- Parallel algorithms
- Scheduling strategies
- Memory aware algorithms
- Sparse tensor computations
- Low rank compression algorithms

Expertise of the team

**Scalable communication optimal algorithms for tensors**

- Analyze existing algorithms
- Determine communication lower bounds
- Propose communication optimal algorithms
- Implement the proposed algorithms

Main focus

**Extension of existing approaches**

- Strassen's concepts to tensors
- Concepts of hierarchical matrices to tensors
- Separation order of dimensions in tensor train

Short/Mid term plans

**Exploratory topics**

- New tensor representations
- Architecture aware algorithms
- Randomization in tensors
- Factorizations of tensors

Mid/Long term plans

## Bringing additional skills in the team

- High dimensional dense tensor computations
- Communication lower bounds for linear algebra computations
- Scalable approaches for large HPC systems
- Use of tensors in quantum and molecular simulations