

Parallel and Scalable Tensor Tools for Modern Computing Systems


Suraj KUMAR

ROMA Team
LIP and Inria, ENS Lyon

AG du LIP
November 29, 2022

Tensors and their uses

- **Neuroscience:** Neuron \times Time \times Trial
 - **Media:** User \times Movie \times Time
 - **Ecommerce:** User \times Product \times Time
 - **Social-Network:** Person \times Person \times Time \times Type
-
- High dimensional tensors: Neural network, Molecular simulation, Quantum computing
 - People work with low dimensional structure (decomposition) of the tensors



Canonical decomposition

$$\text{Tensor} = \text{Column 1} + \text{Column 2} + \cdots + \text{Column N}$$

Tucker decomposition

$$\text{Tensor} = \text{Core Tensor} \times \text{Matrix 1} \times \text{Matrix 2} \times \cdots \times \text{Matrix N}$$

Tensor-train decomposition


$$\text{Tensor} = \text{Column 1} \times \text{Column 2} \times \cdots \times \text{Column N}$$

Importance of communication in high performance computing

- Gaps between computation and communication costs growing exponentially

	Annual improvements
Time-per-operation	59 %
Network-bandwidth	26 %
Network-latency	15 %

Source: Getting up to speed: The future of supercomputing (observed from 2004)



- Research goal:** Scalable and communication optimal tools for tensor computations

Communication Optimal Algorithms for Multiple Tensor-Times-Matrix Computation

Hussam AL DAAS¹, Grey BALLARD², Laura GRIGORI³, Suraj KUMAR⁴, and Kathryn ROUSE⁵

¹Rutherford Appleton Laboratory, UK

²Wake Forest University, USA


³Inria Paris, France

⁴Inria and ENS Lyon, France

⁵Inmar Intelligence, USA

AG du LIP
November 29, 2022

Higher-order SVD (HOSVD) to compute Tucker decomposition



Algorithm 1 3-dimensional HOSVD Algorithm(\mathcal{X})

- 1: Obtain factor matrices A , B and C from the matrix representations of the input tensor \mathcal{X}
 - 2: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$
 - 3: Return \mathcal{Y} , A , B , C
-


- \mathcal{X}, \mathcal{Y} : 3-dimensional input and output tensors (or arrays) & A, B, C : matrices
- \times_i : tensor contraction along the i th dimension (similar to matrix multiplication)
- Multiple Tensor-Times-Matrix (Multi-TTM) computation: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$

Settings

- P number of processors
- Each processor performs (asymptotically) equal amount of operations
- One copy of data is in the system
 - $1/P$ th amount of inputs (before the computation) and output (after the computation) on each processor
- Focus on bandwidth cost (volume of data transfers)

Approach to obtain communication lower bounds

- Loomis-Whitney inequality: for $d - 1$ dimensional projections
 - For the 2d object A , $\phi_x \phi_y \geq \text{Area}(A)$
 - For the 3d object B , $(\phi_{xy} \phi_{yz} \phi_{xz})^{\frac{1}{2}} \geq \text{Volume}(B)$
- Hölder-Brascamp-Lieb (HBL) inequality – generalization for arbitrary dimensional projections
 - Provide exponent for each projection



Constraints for parallel load balanced matrix matrix multiplication

- $C = AB$ with $A \in \mathbb{R}^{n_1 \times n_2}$, $B \in \mathbb{R}^{n_2 \times n_3}$, and $C \in \mathbb{R}^{n_1 \times n_3}$
$$\text{for } i = 1:n_1, \text{ for } k = 1:n_2, \text{ for } j = 1:n_3$$
$$C[i][j] += A[i][k] * B[k][j]$$
- ϕ_A, ϕ_B, ϕ_C : projections of computations on arrays A, B, C
- From Loomis-Whitney/HBL inequality: $\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \text{number of multiplications per processor} = \frac{n_1 n_2 n_3}{P}$
- Extra constraints: $\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$, $\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$, $\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$

Optimization problem and communication lower bounds

Minimize $\phi_A + \phi_B + \phi_C$ s.t.

$$\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \frac{n_1 n_2 n_3}{P}$$


$$\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$$

$$\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$$

$$\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$$

Amount of array accesses = $\phi_A + \phi_B + \phi_C$

- Estimate the solution and prove optimality using all Karush–Kuhn–Tucker conditions are satisfied
- For $n_1 \leq n_2 \leq n_3$,




- Communication lower bound = $\phi_A + \phi_B + \phi_C - \text{data owned by the processor}$ = $\phi_A + \phi_B + \phi_C - \frac{n_1 n_2 + n_2 n_3 + n_1 n_3}{P}$

Design of communication optimal algorithms for $C = AB$

Arrangements of 8 processors




- P is organized into $p_1 \times p_2 \times p_3$ logical grid
- Select p_1, p_2 and p_3 based on the communication lower bounds
- Gather A on the set of processors along each slice of p_3
- Gather B on the set of processors along each slice of p_1
- Perform local computation
- Perform reduce operation along p_2 to obtain C



Outline

3-dimensional Multi-TTM computation



- $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$
- \mathcal{X}, \mathcal{Y} : 3-dimensional input and output tensors
- A, B, C : matrices
- \times_i : analogous to matrix multiplication

- TTM-in-Sequence approach (used in TuckerMPI library): $\mathcal{Y} = ((\mathcal{X} \times_1 A^T) \times_2 B^T) \times_3 C^T$
- Our All-at-Once definition with $\mathcal{X} \in \mathbb{R}^{n_1 \times n_2 \times n_3}$, $\mathcal{Y} \in \mathbb{R}^{r_1 \times r_2 \times r_3}$, $A \in \mathbb{R}^{n_1 \times r_1}$, $B \in \mathbb{R}^{n_2 \times r_2}$, $C \in \mathbb{R}^{n_3 \times r_3}$

for $\{n'_1, n'_2, n'_3, r'_1, r'_2, r'_3\} = 1:\{n_1, n_2, n_3, r_1, r_2, r_3\}$

$$\mathcal{Y}(r'_1, r'_2, r'_3) += (\mathcal{X}(n'_1, n'_2, n'_3) * A(n'_1, r'_1) * B(n'_2, r'_2) * C(n'_3, r'_3))$$

- Establish lower bounds on data transfers based on the geometry of computations
- Design a 6-dimensional parallel algorithm
 - Select right parameters based on lower bounds to achieve communication optimality

Solving optimization problems to compute lower bounds

- Select a processor which performs $\frac{n_1 r_1 n_2 r_2 n_3 r_3}{P}$ amount of 4 – array operations
- After applying lower and upper bounds for each projection, we need to solve the following optimization problem

Minimize $\phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3$ s.t.

$$\phi_x^{1-a} \phi_y^{1-a} \phi_1^a \phi_2^a \phi_3^a \geq \frac{n_1 r_1 n_2 r_2 n_3 r_3}{P}$$

$$\frac{n_1 n_2 n_3}{P} \leq \phi_x \leq n_1 n_2 n_3$$

$$\frac{r_1 r_2 r_3}{P} \leq \phi_y \leq r_1 r_2 r_3$$

$$\frac{n_1 r_1}{P} \leq \phi_1 \leq n_1 r_1$$

$$\frac{n_2 r_2}{P} \leq \phi_2 \leq n_2 r_2$$

$$\frac{n_3 r_3}{P} \leq \phi_3 \leq n_3 r_3$$

$$0 \leq a \leq 1$$

Divide the problem into two parts

Matrix part

Minimize $\phi_1 + \phi_2 + \phi_3$ s.t.

$$\phi_1\phi_2\phi_3 \geq \frac{n_1r_1n_2r_2n_3r_3}{P}$$

$$\frac{n_1r_1}{P} \leq \phi_1 \leq n_1r_1$$

$$\frac{n_2r_2}{P} \leq \phi_2 \leq n_2r_2$$

$$\frac{n_3r_3}{P} \leq \phi_3 \leq n_3r_3$$

Tensor part

Minimize $\phi_x + \phi_y$ s.t.

$$\phi_x\phi_y \geq \frac{n_1r_1n_2r_2n_3r_3}{P}$$


$$\frac{n_1n_2n_3}{P} \leq \phi_x \leq n_1n_2n_3$$

$$\frac{r_1r_2r_3}{P} \leq \phi_y \leq r_1r_2r_3$$

Amount of accesses and lower bounds

- We assume $n_1 r_1 \leq n_2 r_2 \leq n_3 r_3$ and $r_1 r_2 r_3 \leq n_1 n_2 n_3$
- Estimate solutions for both parts using Lagrange multipliers (optimality can be proven using Karush–Kuhn–Tucker conditions)

Amount of accesses = $\phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3$




Communication lower bound = $\phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3 - \frac{n_1 n_2 n_3 + r_1 r_2 r_3 + n_1 r_1 + n_2 r_2 + n_3 r_3}{P}$

Design of communication optimal algorithms

Data Distribution (P is organized into a $p_1 \times p_2 \times p_3 \times q_1 \times q_2 \times q_3$ grid)

- p_1, p_2, p_3, q_1, q_2 , and q_3 evenly distribute n_1, n_2, n_3, r_1, r_2 , and r_3
- Each processor has $\frac{1}{P}$ th amount of input and output variables
- Subtensor $\mathcal{X}_{231} = \mathcal{X}\left(\frac{n_1}{p_1} + 1 : 2\frac{n_1}{p_1}, 2\frac{n_2}{p_2} + 1 : 3\frac{n_2}{p_2}, 1 : \frac{n_3}{p_3}\right)$ is distributed evenly among processors $(2, 3, 1, *, *, *)$
- Submatrix $B_{31} = B\left(2\frac{n_2}{p_2} + 1 : 3\frac{n_2}{p_2}, 1 : \frac{r_2}{q_2}\right)$ is distributed evenly among processors $(*, 3, *, *, 1, *)$



6-dimensional algorithm to compute Multi-TTM

Algorithm 1 3-dimensional Parallel Atomic Multi-TTM

Require: $\mathcal{X}, A, B, C, p_1 \times p_2 \times p_3 \times q_1 \times q_2 \times q_3$ logical processor grid


Ensure: \mathcal{Y} such that $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$

- 1: $(p'_1, p'_2, p'_3, q'_1, q'_2, q'_3)$ is my processor id
 - 2: //All-gather input tensor and matrices
 - 3: $\mathcal{X}_{p'_1 p'_2 p'_3} = \text{All-Gather}(\mathcal{X}, (p'_1, p'_2, p'_3, *, *, *))$
 - 4: $A_{p'_1 q'_1} = \text{All-Gather}(A, (p'_1, *, *, q'_1, *, *))$
 - 5: $B_{p'_2 q'_2} = \text{All-Gather}(B, (*, p'_2, *, *, q'_2, *))$
 - 6: $C_{p'_3 q'_3} = \text{All-Gather}(C, (*, *, p'_3, *, *, q'_3))$
 - 7: //Perform local Multi-TTM computation in a temporary tensor \mathcal{T}
 - 8: $\mathcal{T} = \text{Local-Multi-TTM}(\mathcal{X}_{p'_1 p'_2 p'_3}, A_{p'_1 q'_1}, B_{p'_2 q'_2}, C_{p'_3 q'_3})$
 - 9: //Reduce-scatter the output tensor in $\mathcal{Y}_{q'_1 q'_2 q'_3}$
 - 10: $\text{Reduce-Scatter}(\mathcal{Y}_{q'_1 q'_2 q'_3}, \mathcal{T}, (*, *, *, q'_1, q'_2, q'_3))$
-


Outline

Communication lower bound (C_{LB}) distributions

$$n_1 = n_2 = n_3 = 2^{12}, r_1 = r_2 = r_3 = 2^4$$




$$n_1 = n_2 = n_3 = 2^{20}, r_1 = r_2 = r_3 = 2^8$$




- Matrix communication costs dominate when P is much less than $\frac{n_1 n_2 n_3}{r_1 r_2 r_3}$

Performance comparison of our algorithm

$$n_1 = n_2 = n_3 = 2^{20}, r_1 = r_2 = r_3 = 2^8$$



$$n_1 = n_2 = n_3 = 2^{12}, r_1 = r_2 = r_3 = 2^4$$



- Typical scenarios in data compression problems
- For small P , our approach communicates much less than the state-of-the-art approach

Outline

Conclusion and future work


Conclusion

- Communication lower bounds and optimal algorithms for All-at-Once Multi-TTM
- Comparison of our approach with the TTM-in-Sequence approach
- Our algorithm communicates much less data than TTM-in-Sequence for small P

Future Work

- Detailed study of what scenarios are favorable for our approach
- Combine both All-at-Once and TTM-in-Sequence approaches
- Extend our framework for other linear algebra computations

Thank You!



=

