

Communication Lower Bounds for Parallel Matrix and Tensor Computations

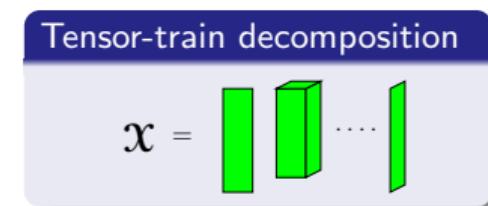
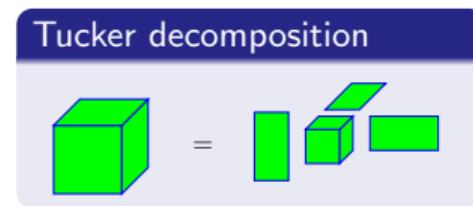
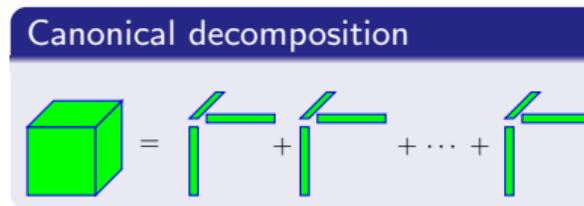
Suraj KUMAR

National Institute for Research in Digital Science and Technology (Inria), Lyon

CSE – IIT Kanpur
May 20, 2025

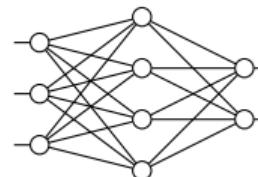
Tensors and their uses

- **Neuroscience:** Neuron \times Time \times Trial
- **Media:** User \times Movie \times Time
- **Ecommerce:** User \times Product \times Time
- **Social-Network:** Person \times Person \times Time \times Type
- High dimensional tensors: Neural network, Molecular simulation, Quantum computing
- People work with low dimensional structure (decomposition) of the tensors



High dimensional tensors: tensorized neural networks

- Neural Network



- Tensor representation of a neural network

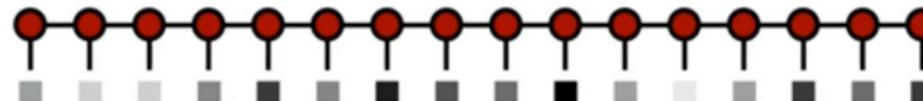


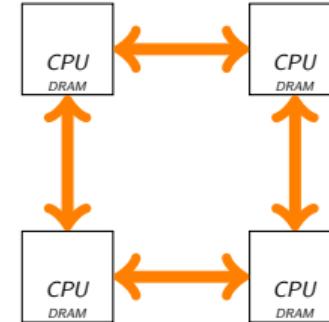
Image source: Perspective on Tensor Networks for Machine Learning - E.M. Stoudenmire

Importance of communication in high performance computing

- Gaps between computation and communication costs growing exponentially

	Annual improvements
Time-per-operation	59 %
Network-bandwidth	26 %
Network-latency	15 %

Source: Getting up to speed: The future of supercomputing (observed from 2004)



- Research goal:** Scalable and communication optimal tools for matrix and tensor computations

Communication Lower Bounds for Parallel Matrix and Tensor Computations

Hussam AL DAAS¹, Grey BALLARD², Laura GRIGORI³, Suraj KUMAR⁴, and Kathryn ROUSE⁵

¹Rutherford Appleton Laboratory, UK

²Wake Forest University, USA

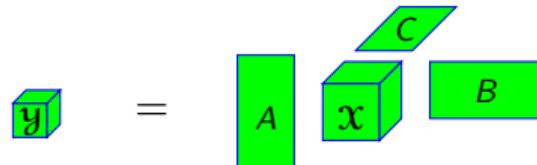
³Ecole Polytechnique Fédérale de Lausanne (EPFL), Switzerland

⁴Inria Lyon, France

⁵Inmar Intelligence, USA

CSE – IIT Kanpur
May 20, 2025

Higher-order SVD (HOSVD) to compute Tucker decomposition



Algorithm 1 3-dimensional HOSVD Algorithm(\mathcal{X})

- 1: Obtain factor matrices A , B and C from the matrix representations of the input tensor \mathcal{X}
 - 2: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$
 - 3: Return \mathcal{Y} , A , B , C
-

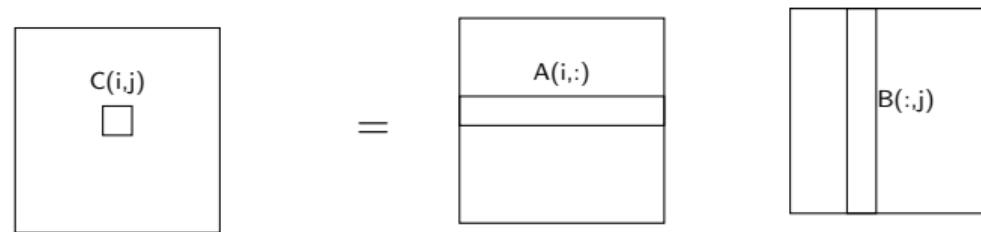
- \mathcal{X} , \mathcal{Y} : 3-dimensional input and output tensors (or arrays) & A , B , C : matrices
- \times_i : tensor contraction along the i th dimension (similar to matrix multiplication)
- Multiple Tensor-Times-Matrix (Multi-TTM) computation: $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$

- 1 For Matrix Matrix Multiplication
- 2 For Multi-TTM Computation

Traditional matrix multiplication

- $C = AB$, where $A \in \mathbb{R}^{n_1 \times n_2}$, $B \in \mathbb{R}^{n_2 \times n_3}$, and $C \in \mathbb{R}^{n_1 \times n_3}$.
- $C_{ij} = \sum_{\ell} A_{i\ell} \cdot B_{\ell j}$

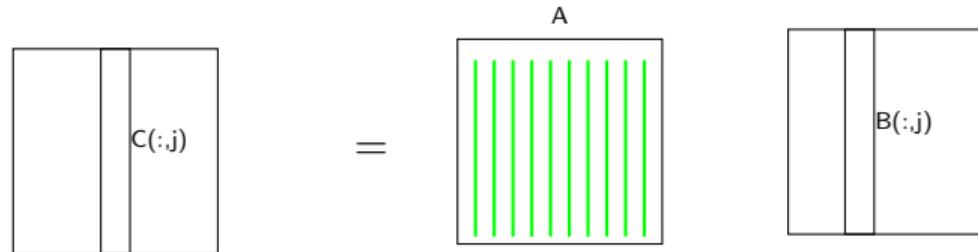
For simplicity, we assume $n_1 = n_2 = n_3 = n$.



- Not focus on fast matrix multiplications such as Strassen's algorithm today

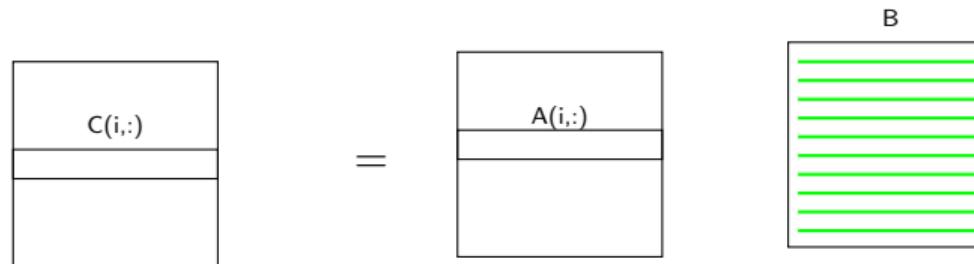
Matrix multiplication: linear combination of columns

- A column of C is obtained by linear combination of columns of A .



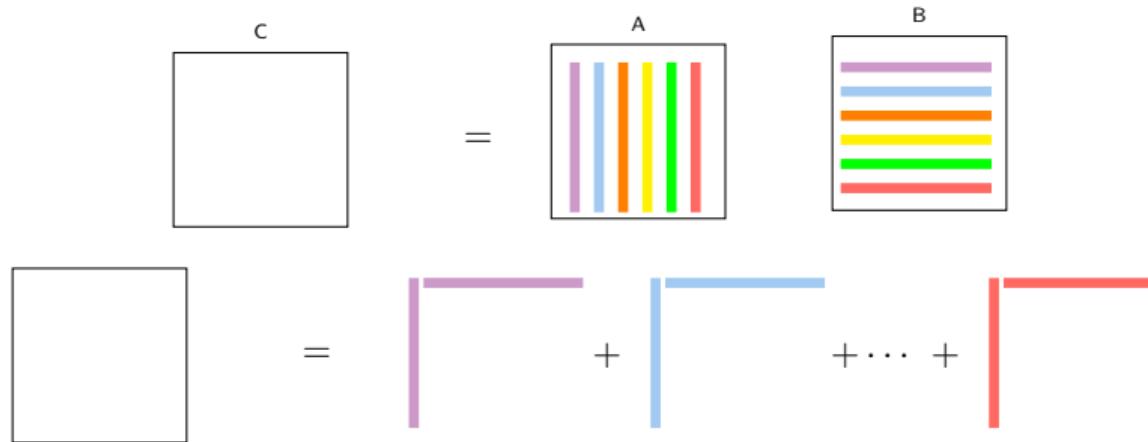
Matrix multiplication: linear combination of rows

- A row of C is obtained by linear combination of rows of B .



Matrix multiplication: sum of n matrices

- Matrix multiplication can also be viewed as sum of n matrices.



Matrix multiplication: recursive calls on submatrices

- Matrix is divided into 2×2 blocks

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}$$

$$C_{11} = A_{11}B_{11} + A_{12}B_{21}$$

$$C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21}$$

$$C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Matrix multiplication: recursive calls on submatrices

Operation count recurrence,

$$T(n) = 8T\left(\frac{n}{2}\right) + \mathcal{O}(n^2)$$

$$T(n) = 1$$

Here $\mathcal{O}(n^2)$ refers that $\exists c \in \mathbb{N}$ such that this term is less than or equal to cn^2 for every n .

After solving, we obtain $T(n) = \mathcal{O}(n^3)$.

Communication lower bounds and communication optimal algorithms

1 For Matrix Matrix Multiplication

- Communication lower bounds & optimal algorithms
- Conclusion

2 For Multi-TTM Computation

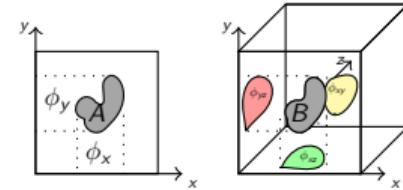
- Communication lower bounds & optimal algorithms
- Simulated experiments
- Conclusion

Settings

- P number of processors
- Each processor performs (asymptotically) equal amount of operations
- One copy of data is in the system
 - $1/P$ th amount of inputs (before the computation) and output (after the computation) on each processor
- Focus on bandwidth cost (volume of data transfers)

Approach to obtain communication lower bounds

- Loomis-Whitney inequality: for $d - 1$ dimensional projections
 - For the 2d object A , $\phi_x \phi_y \geq \text{Area}(A)$
 - For the 3d object B , $(\phi_{xy} \phi_{yz} \phi_{xz})^{\frac{1}{2}} \geq \text{Volume}(B)$
- Hölder-Brascamp-Lieb (HBL) inequality – generalization for arbitrary dimensional projections
 - Provide exponent for each projection



Constraints for parallel load balanced matrix matrix multiplication

- $C = AB$ with $A \in \mathbb{R}^{n_1 \times n_2}$, $B \in \mathbb{R}^{n_2 \times n_3}$, and $C \in \mathbb{R}^{n_1 \times n_3}$
 - for $i = 1:n_1$, for $k = 1:n_2$, for $j = 1:n_3$
$$C[i][j] += A[i][k] * B[k][j]$$
- ϕ_A, ϕ_B, ϕ_C : projections of computations on arrays A, B, C
- From Loomis-Whitney/HBL inequality: $\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \text{number of multiplications per processor} = \frac{n_1 n_2 n_3}{P}$
- Extra constraints: $\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$, $\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$, $\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$

Optimization problem and communication lower bounds

Minimize $\phi_A + \phi_B + \phi_C$ s.t.

$$\phi_A^{\frac{1}{2}} \phi_B^{\frac{1}{2}} \phi_C^{\frac{1}{2}} \geq \frac{n_1 n_2 n_3}{P}$$

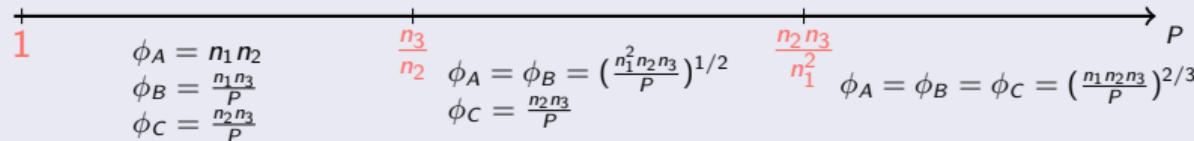
$$\frac{n_1 n_2}{P} \leq \phi_A \leq n_1 n_2$$

$$\frac{n_2 n_3}{P} \leq \phi_B \leq n_2 n_3$$

$$\frac{n_1 n_3}{P} \leq \phi_C \leq n_1 n_3$$

Amount of array accesses = $\phi_A + \phi_B + \phi_C$

- Estimate the solution and prove optimality using all Karush–Kuhn–Tucker conditions are satisfied
- For $n_1 \leq n_2 \leq n_3$,



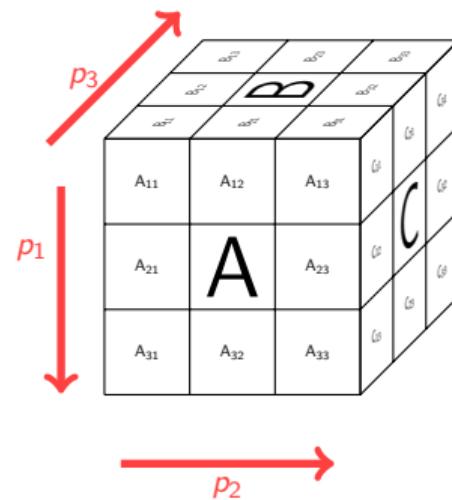
- Communication lower bound = $\phi_A + \phi_B + \phi_C - \text{data owned by the processor} = \phi_A + \phi_B + \phi_C - \frac{n_1 n_2 + n_2 n_3 + n_1 n_3}{P}$

Design of communication optimal algorithms for $C = AB$

Arrangements of 8 processors



- P is organized into $p_1 \times p_2 \times p_3$ logical grid
- Select p_1, p_2 and p_3 based on the communication lower bounds
- Gather A on the set of processors along each slice of p_3
- Gather B on the set of processors along each slice of p_1
- Perform local computation
- Perform reduce operation along p_2 to obtain C



Outline

1 For Matrix Matrix Multiplication

- Communication lower bounds & optimal algorithms
- Conclusion

2 For Multi-TTM Computation

- Communication lower bounds & optimal algorithms
- Simulated experiments
- Conclusion

Conclusion and future work

Conclusion

- Communication lower bounds for matrix multiplication
- Communication optimal algorithms for matrix multiplication

Future Work

- Obtain tight communication lower bounds for Strassen's algorithm
- Analyze symmetric computations such as $C = AA^T$ in detail

Outline

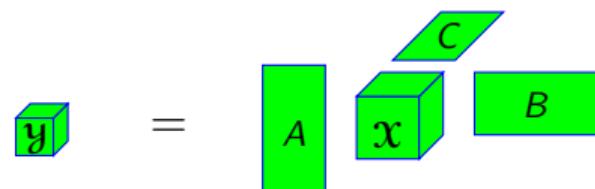
1 For Matrix Matrix Multiplication

- Communication lower bounds & optimal algorithms
- Conclusion

2 For Multi-TTM Computation

- Communication lower bounds & optimal algorithms
- Simulated experiments
- Conclusion

3-dimensional Multi-TTM computation



- $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$
- \mathcal{X}, \mathcal{Y} : 3-dimensional input and output tensors
- A, B, C : matrices
- \times_i : analogous to matrix multiplication

- TTM-in-Sequence approach (used in TuckerMPI library): $\mathcal{Y} = ((\mathcal{X} \times_1 A^T) \times_2 B^T) \times_3 C^T$
- Our All-at-Once definition with $\mathcal{X} \in \mathbb{R}^{n_1 \times n_2 \times n_3}$, $\mathcal{Y} \in \mathbb{R}^{r_1 \times r_2 \times r_3}$, $A \in \mathbb{R}^{n_1 \times r_1}$, $B \in \mathbb{R}^{n_2 \times r_2}$, $C \in \mathbb{R}^{n_3 \times r_3}$

for $\{n'_1, n'_2, n'_3, r'_1, r'_2, r'_3\} = 1:\{n_1, n_2, n_3, r_1, r_2, r_3\}$

$$\mathcal{Y}(r'_1, r'_2, r'_3) += (\mathcal{X}(n'_1, n'_2, n'_3) * A(n'_1, r'_1) * B(n'_2, r'_2) * C(n'_3, r'_3))$$

- Establish lower bounds on data transfers based on the geometry of computations
- Design a 6-dimensional parallel algorithm
 - Select right parameters based on lower bounds to achieve communication optimality

Solving optimization problems to compute lower bounds

- Select a processor which performs $\frac{n_1 r_1 n_2 r_2 n_3 r_3}{P}$ amount of 4 – array operations
- After applying lower and upper bounds for each projection, we need to solve the following optimization problem

$$\text{Minimize } \phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3 \text{ s.t.}$$

$$\phi_x^{1-a} \phi_y^{1-a} \phi_1^a \phi_2^a \phi_3^a \geq \frac{n_1 r_1 n_2 r_2 n_3 r_3}{P}$$

$$\frac{n_1 n_2 n_3}{P} \leq \phi_x \leq n_1 n_2 n_3$$

$$\frac{r_1 r_2 r_3}{P} \leq \phi_y \leq r_1 r_2 r_3$$

$$\frac{n_1 r_1}{P} \leq \phi_1 \leq n_1 r_1$$

$$\frac{n_2 r_2}{P} \leq \phi_2 \leq n_2 r_2$$

$$\frac{n_3 r_3}{P} \leq \phi_3 \leq n_3 r_3$$

$$0 \leq a \leq 1$$

Divide the problem into two parts

Matrix part

Minimize $\phi_1 + \phi_2 + \phi_3$ s.t.

$$\phi_1\phi_2\phi_3 \geq \frac{n_1r_1n_2r_2n_3r_3}{P}$$

$$\frac{n_1r_1}{P} \leq \phi_1 \leq n_1r_1$$

$$\frac{n_2r_2}{P} \leq \phi_2 \leq n_2r_2$$

$$\frac{n_3r_3}{P} \leq \phi_3 \leq n_3r_3$$

Tensor part

Minimize $\phi_x + \phi_y$ s.t.

$$\phi_x\phi_y \geq \frac{n_1r_1n_2r_2n_3r_3}{P}$$

$$\frac{n_1n_2n_3}{P} \leq \phi_x \leq n_1n_2n_3$$

$$\frac{r_1r_2r_3}{P} \leq \phi_y \leq r_1r_2r_3$$

Amount of accesses and lower bounds

- We assume $n_1 r_1 \leq n_2 r_2 \leq n_3 r_3$ and $r_1 r_2 r_3 \leq n_1 n_2 n_3$
- Estimate solutions for both parts using Lagrange multipliers (optimality can be proven using Karush–Kuhn–Tucker conditions)

Amount of accesses = $\phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3$



$$\phi_1 = n_1 r_1$$

$$\phi_2 = n_2 r_2$$

$$\phi_3 = \frac{n_3 r_3}{P}$$

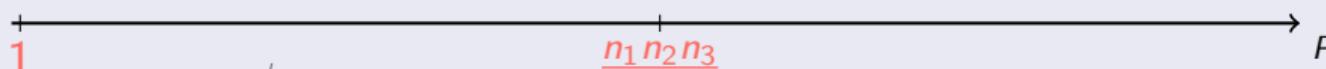
$$\frac{n_3 r_3}{n_2 r_2}$$

$$\phi_1 = n_1 r_1$$

$$\phi_2 = \phi_3 = \left(\frac{n_2 r_2 n_3 r_3}{P}\right)^{1/2}$$

$$\frac{n_2 r_2 n_3 r_3}{n_1^2 r_1^2}$$

$$\phi_1 = \phi_2 = \phi_3 = \left(\frac{n_1 r_1 n_2 r_2 n_3 r_3}{P}\right)^{1/3}$$



$$\phi_y = r_1 r_2 r_3$$

$$\phi_x = \frac{n_1 n_2 n_3}{P}$$

$$\frac{n_1 n_2 n_3}{r_1 r_2 r_3}$$

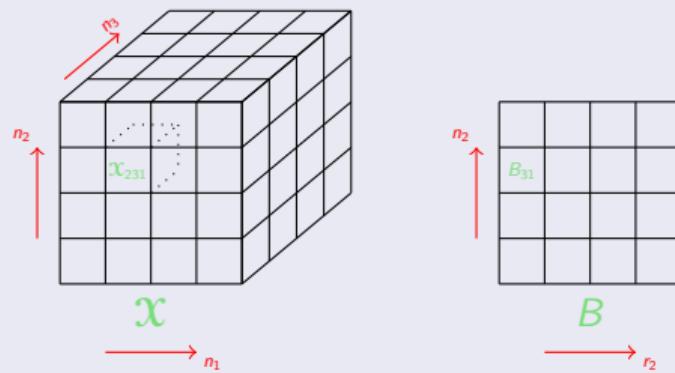
$$\phi_x = \phi_y = \left(\frac{n_1 n_2 n_3 r_1 r_2 r_3}{P}\right)^{1/2}$$

Communication lower bound = $\phi_x + \phi_y + \phi_1 + \phi_2 + \phi_3 - \frac{n_1 n_2 n_3 + r_1 r_2 r_3 + n_1 r_1 + n_2 r_2 + n_3 r_3}{P}$

Design of communication optimal algorithms

Data Distribution (P is organized into a $p_1 \times p_2 \times p_3 \times q_1 \times q_2 \times q_3$ grid)

- p_1, p_2, p_3, q_1, q_2 , and q_3 evenly distribute n_1, n_2, n_3, r_1, r_2 , and r_3
- Each processor has $\frac{1}{P}$ th amount of input and output variables
- Subtensor $\mathcal{X}_{231} = \mathcal{X}\left(\frac{n_1}{p_1} + 1 : 2\frac{n_1}{p_1}, 2\frac{n_2}{p_2} + 1 : 3\frac{n_2}{p_2}, 1 : \frac{n_3}{p_3}\right)$ is distributed evenly among processors $(2, 3, 1, *, *, *)$
- Submatrix $B_{31} = B\left(2\frac{n_2}{p_2} + 1 : 3\frac{n_2}{p_2}, 1 : \frac{r_2}{q_2}\right)$ is distributed evenly among processors $(*, 3, *, *, 1, *)$



6-dimensional algorithm to compute Multi-TTM

Algorithm 1 3-dimensional Parallel Atomic Multi-TTM

Require: $\mathcal{X}, A, B, C, p_1 \times p_2 \times p_3 \times q_1 \times q_2 \times q_3$ logical processor grid

Ensure: \mathcal{Y} such that $\mathcal{Y} = \mathcal{X} \times_1 A^T \times_2 B^T \times_3 C^T$

- 1: $(p'_1, p'_2, p'_3, q'_1, q'_2, q'_3)$ is my processor id
 - 2: //All-gather input tensor and matrices
 - 3: $\mathcal{X}_{p'_1 p'_2 p'_3} = \text{All-Gather}(\mathcal{X}, (p'_1, p'_2, p'_3, *, *, *))$
 - 4: $A_{p'_1 q'_1} = \text{All-Gather}(A, (p'_1, *, *, q'_1, *, *))$
 - 5: $B_{p'_2 q'_2} = \text{All-Gather}(B, (*, p'_2, *, *, q'_2, *))$
 - 6: $C_{p'_3 q'_3} = \text{All-Gather}(C, (*, *, p'_3, *, *, q'_3))$
 - 7: //Perform local Multi-TTM computation in a temporary tensor \mathcal{T}
 - 8: $\mathcal{T} = \text{Local-Multi-TTM}(\mathcal{X}_{p'_1 p'_2 p'_3}, A_{p'_1 q'_1}, B_{p'_2 q'_2}, C_{p'_3 q'_3})$
 - 9: //Reduce-scatter the output tensor in $\mathcal{Y}_{q'_1 q'_2 q'_3}$
 - 10: $\text{Reduce-Scatter}(\mathcal{Y}_{q'_1 q'_2 q'_3}, \mathcal{T}, (*, *, *, q'_1, q'_2, q'_3))$
-

Outline

1 For Matrix Matrix Multiplication

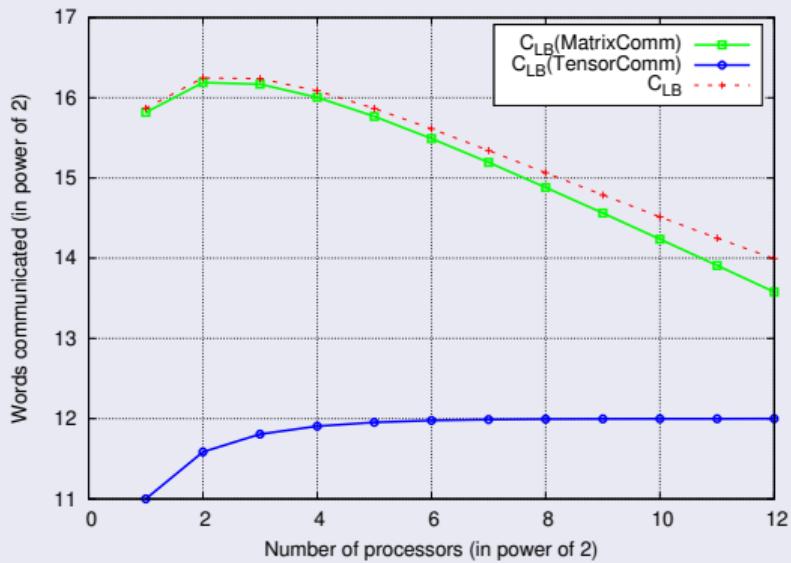
- Communication lower bounds & optimal algorithms
- Conclusion

2 For Multi-TTM Computation

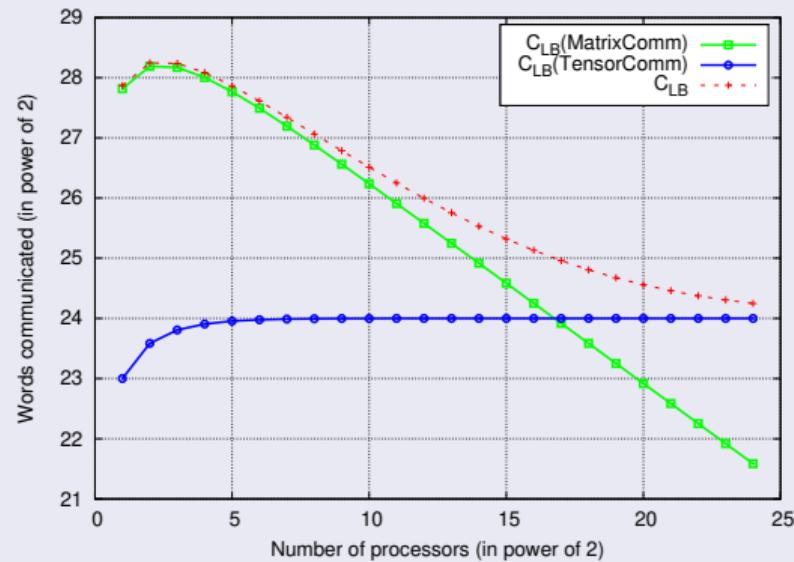
- Communication lower bounds & optimal algorithms
- Simulated experiments
- Conclusion

Communication lower bound (C_{LB}) distributions

$$n_1 = n_2 = n_3 = 2^{12}, r_1 = r_2 = r_3 = 2^4$$



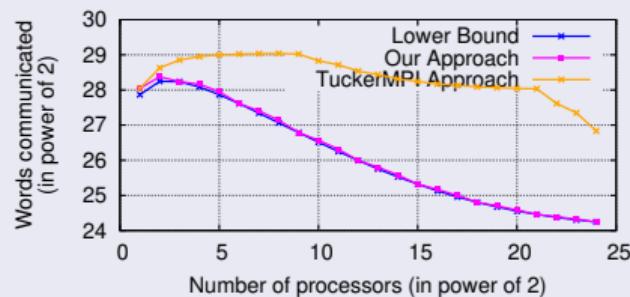
$$n_1 = n_2 = n_3 = 2^{20}, r_1 = r_2 = r_3 = 2^8$$



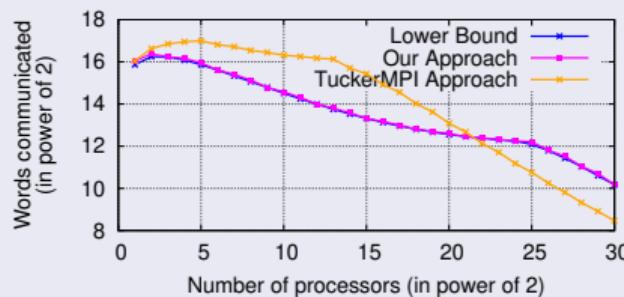
- Matrix communication costs dominate when P is much less than $\frac{n_1 n_2 n_3}{r_1 r_2 r_3}$

Performance comparison of our algorithm

$$n_1 = n_2 = n_3 = 2^{20}, r_1 = r_2 = r_3 = 2^8$$



$$n_1 = n_2 = n_3 = 2^{12}, r_1 = r_2 = r_3 = 2^4$$



- Typical scenarios in data compression problems
- For small P , our approach communicates much less than the state-of-the-art approach

Outline

1 For Matrix Matrix Multiplication

- Communication lower bounds & optimal algorithms
- Conclusion

2 For Multi-TTM Computation

- Communication lower bounds & optimal algorithms
- Simulated experiments
- Conclusion

Conclusion and future work

Conclusion

- Communication lower bounds and optimal algorithms for All-at-Once Multi-TTM
- Comparison of our approach with the TTM-in-Sequence approach
- Our algorithm communicates much less data than TTM-in-Sequence for small P

Future Work

- Detailed study of what scenarios are favorable for our approach
- Combine both All-at-Once and TTM-in-Sequence approaches
- Extend our framework for other linear algebra computations

Thank You!